

Artificial Intelligence

Blai Bonet

Universidad Simón Bolívar, Caracas, Venezuela



Constraint satisfaction problems (CSPs)

© 2018 Blai Bonet

Goals for the lecture

- Constraint satisfaction problem (CSP)
- Types of CSPs and constraints
- Translation of CSPs
- Backtracking algorithms with heuristics for variable selection
- Inference: forward checking, arc consistency
- Solving CSPs by pure inference

© 2018 Blai Bonet

Informal description

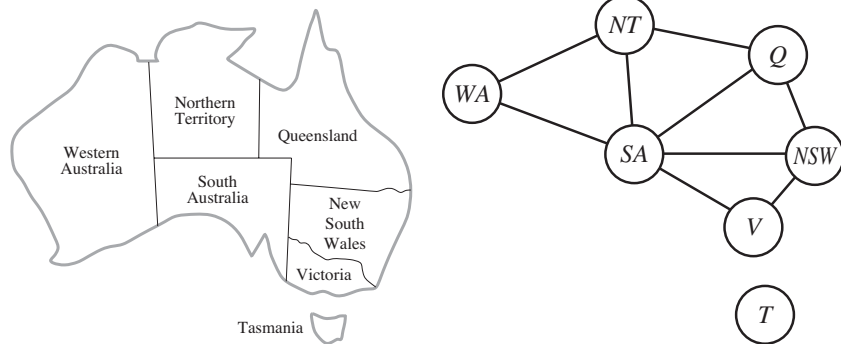
CSP is **assignment problem** defined by:

- a set of **variables** with **domains**
- a set of **constraints**

Task: find assignment of variables to values that **satisfies** all constraints

© 2018 Blai Bonet

Example: Map coloring



[Image from Russell & Norvig. *Artificial Intelligence: A Modern Approach*]

Formal model

CSP is given by tuple $(\mathcal{X}, \mathcal{D}, \mathcal{C})$ where:

- $\mathcal{X} = \{X_1, \dots, X_n\}$ is finite set of variables
- $\mathcal{D} = \{D_1, \dots, D_n\}$ is set of domains, domain D_i for variable X_i
- $\mathcal{C} = \{C_1, \dots, C_m\}$ is set of constraints that specify **allowable combinations** of values

Each constraint C is pair $\langle \text{scope}, R \rangle$ where scope is tuple over X that specifies the variables involved in C , and $R \subseteq \prod_{X_i \in \text{scope}} D_i$ defines the allowable combinations for variables in scope

E.g., if X and Y are binary variables with domain $\{0, 1\}$, the constraint $\langle (X, Y), \{(0, 1), (1, 0)\} \rangle$ expresses $X \neq Y$

Formulation of map coloring

- Variables $\mathcal{X} = \{WA, NT, Q, NSW, V, SA, T\}$
- Domains for all variables given by colors *red*, *blue* and *green*
- For each two variables X and Y **connected by edge**, there is a constraint with scope (X, Y) and the following relation that requires that the colors of X and Y must be different:

$\{(red, blue), (red, green), (blue, red), (blue, green), (green, red), (green, blue)\}$

Solving CSPs by search

CSPs can be solved by performing search on the space of **partial assignments** of variables to values:

- initial state is **empty assignment**
- goal states are **complete assignments** that satisfy all constraints
- successor function **extends** partial assignment with new variable, provided that resulting assignment is consistent (i.e. doesn't violate a constraint)
- uniform costs

If there are n variables, all goal states (if any) appear at depth n

IDA* is discarded. DFBnB could be considered but there are no **meaningful heuristics** since all costs are equal. We'll do a depth-first traversal but extended with some form of "inference" to **prune branches in search tree**

Alternative model for (local) search

Another search space is obtained by considering only **complete assignments** instead of partial ones

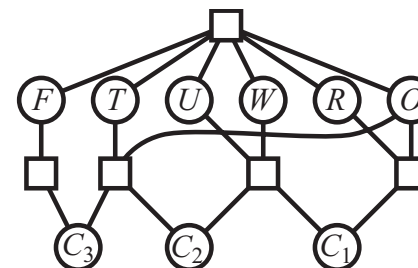
Edges connect assignment that differ in the assignment for one or more variables (typically just 1 variable)

Initial state is any assignment while goal states correspond to assignments that satisfy all constraints

Formulation used by **local search methods** that in some cases are very effective but **incomplete**

Example: Cryptoarithmic

$$\begin{array}{r} T W O \\ + T W O \\ \hline F O U R \end{array}$$

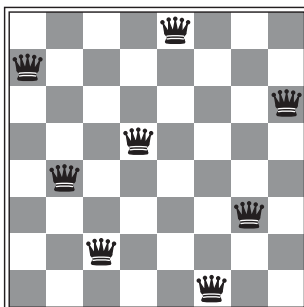


[Image from Russell & Norvig. *Artificial Intelligence: A Modern Approach*]

Example: 8-Queens

Place 8 queens in an empty chess board in a way that no queen attacks another

Can it be done?



[Image from Russell & Norvig. *Artificial Intelligence: A Modern Approach*]

CSP Variations

Simplest CSPs have **finite and discrete domains**

Infinite discrete domains can be considered, but constraints cannot be represented explicitly and **constraint languages** are used

Continuous domains such as real values can also be considered

Some special cases:

- Real-valued variables with **linear constraints** (e.g. $X + 3Y \leq Z$) can be solved efficiently with **linear programming**
- Integer-valued variables with linear constraints can be solved using **integer programming** methods (intractable in worst case)
- Special cases like real-valued variables with **convex constraints**

Constraint types and constraint graph

A constraint whose scope is singleton is **unary constraint**

A **binary constraint** relates two variables (scope size is 2)

A **constraint of order k** relates k variables; for $k > 2$, it is a **higher-order** constraint

Constraint graph for CSP $(\mathcal{X}, \mathcal{D}, \mathcal{C})$ is **(undirected) graph** with vertices given by \mathcal{X} and edges (X_i, X_j) iff there is a constraint whose scope contains i and j

CSP with binary constraints

Any CSP $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ can be mapped into **equivalent** CSP $P' = (\mathcal{X}', \mathcal{D}', \mathcal{C}')$ with $\mathcal{X}' \supseteq \mathcal{X}$ (i.e. with possibly more variables) but with **binary constraints**

Equivalent means:

- any solution for P can be extended into a solution for P'
- any solution for P' corresponds to a solution for P (i.e. if ν is a solution for P' , then its projection $\nu|_{\mathcal{X}}$ over \mathcal{X} is a solution for P)

Mapping CSPs to binary CSPs

For $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ with n vars and m constraints, define $P' = (\mathcal{X}', \mathcal{D}', \mathcal{C}')$:

- $\mathcal{X}' = \{X'_i : X_i \in \mathcal{X}\} \cup \{X_{n+j} : 1 \leq j \leq m\}$ (one **new var** per constr.)
- Domains for original vars: $D'_i = D_i \cap \cap \{R_j : \text{scope}_j = (X_i)\}$, $i = 1 \dots n$
- Domains for new vars: $D'_{n+j} = R_j$, $j = 1 \dots m$ (var X'_{n+j} has domain given by tuples permitted by constraint C_j : $D'_{n+j} \subseteq \Pi_{X_i \in \text{scope}_j} D_i$)
- **Binary constraints**: for each (i, j) such that $X_i \in \text{scope}_j$, add constraint $C'_{i,j} = \langle \text{scope}'_{i,j}, R'_{i,j} \rangle$ where:
 - $\text{scope}'_{i,j} = (X'_i, X'_{n+j})$
 - $R'_{i,j} = \{(x_i, t) \in D'_i \times D'_{n+j} : t[X_i] = x_i\}$

Example: Mapping CSP to binary CSP

Problem P with variables $\mathcal{X} = \{X_1, X_2, X_3\}$ over domain $D = \{0, 1, 2\}$ and two constraints: $X_3 = X_1 + X_2 \bmod 3$, and $X_2 + X_3 \geq 1$

Transformed problem is $P' = (\mathcal{X}' = \{X'_i : 1 \leq i \leq 5\}, \mathcal{D}', \mathcal{C}')$ where

- $D'_i = D$ for $i = 1, 2, 3$
- $D'_4 = \{(x_1, x_2, x_3) \in D^3 : x_3 = x_1 + x_2 \bmod 3\}$
- $D'_5 = \{(x_2, x_3) \in D^2 : x_2 + x_3 \geq 1\}$
- $C'_{14} = \langle (X'_1, X'_4), R'_{14} \rangle$ with $R'_{14} = \{(x_1, (x_1, x_2, x_3)) : x_1 \in D, (x_1, x_2, x_3) \in D'_4\}$
- $C'_{24} = \langle (X'_2, X'_4), R'_{24} \rangle$ with $R'_{24} = \{(x_2, (x_1, x_2, x_3)) : x_2 \in D, (x_1, x_2, x_3) \in D'_4\}$
- $C'_{34} = \langle (X'_3, X'_4), R'_{34} \rangle$ with $R'_{34} = \{(x_3, (x_1, x_2, x_3)) : x_3 \in D, (x_1, x_2, x_3) \in D'_4\}$
- $C'_{25} = \langle (X'_2, X'_5), R'_{25} \rangle$ with $R'_{25} = \{(x_2, (x_2, x_3)) : x_2 \in D, (x_2, x_3) \in D'_5\}$
- $C'_{35} = \langle (X'_3, X'_5), R'_{35} \rangle$ with $R'_{35} = \{(x_3, (x_2, x_3)) : x_3 \in D, (x_2, x_3) \in D'_5\}$

Example: Sudoku

	1	2	3	4	5	6	7	8	9
A			3		2		6		
B	9			3		5			1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8			2	3				9
I			5		1		3		

[Image from Russell & Norvig. *Artificial Intelligence: A Modern Approach*]

Naive backtracking algorithm

For given node n , children of n correspond to different **extensions with one variable** of the assignment associated to n

```

1 naive-backtrack(assignment A, csp P)
2   if A is complete assignment then return A
3
4   foreach variable X unassigned by A
5     foreach value in domain of X
6       if X = value is consistent with A wrt P
7         A' := A union { X = value }
8         result := naive-backtrack(A', P)
9         if result != FAIL then return result
10
11  return FAIL

```

For n variables and $d = \max_i |D_i|$, branching factor at root is $O(nd)$, at second level $O((n-1)d)$, etc. Total number of leaves is $O(n!d^n)$, yet number of assignments is only $O(d^n)$

Basic backtracking algorithm

For given node n , children of n correspond to the different values for **fixed unassigned variable**

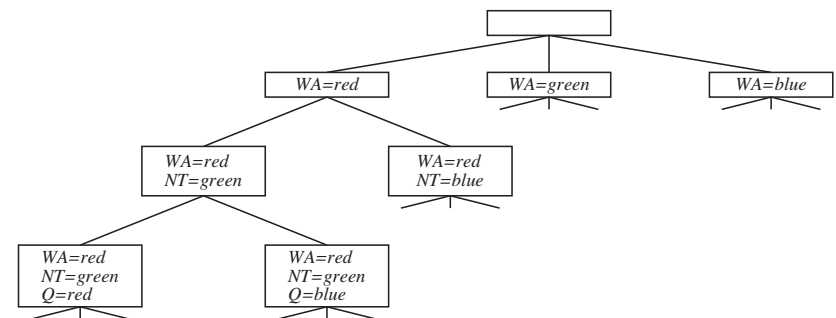
```

1 backtrack(assignment A, csp P)
2   if A is complete assignment then return A
3
4   X := select-unassigned-variable(A, P)
5   foreach value in domain of X
6     if X = value is consistent with A wrt P
7       A' := A union { X = value }
8       result := backtrack(A', P)
9       if result != FAIL then return result
10
11  return FAIL

```

Branching factor is $O(d)$, where $d = \max_i |D_i|$. With n variables, number of leaves is $O(d^n)$ and equal to number of assignments

Example: Backtracking



[Image from Russell & Norvig. *Artificial Intelligence: A Modern Approach*]

Critical issues when implementing solution

- Which variable should be chosen at each node? How should its values be ordered for the recursion?
- What are the implications of current assignment for still unassigned variables?
- When branch fails, can the search avoid repeating the failure in next branches?

Variable ordering

Idea is to choose the **most constrained variable** in order to detect a failure (backtrack) as soon as possible

It is better to fail high on a branch than deep. Heuristic is called **MRV (Minimum Remaining Values)**, Most Constrained Variable, or “fail-first”

Another idea is to choose variable involved in most constraints. It can be combined with MRV as a **tie-breaker**:

If two variables have the same number of remaining values (MRV criterion), prefer the one involved in more constraints

Value ordering

Once a variable is selected, its values must be ordered

Least-constraining value is an effective heuristic:

*Prefer values that **rule out fewest values** for neighbor variables in constraint graph*

Motivation is that once a variable is fixed, the algorithm should try to find a solution as fast as possible

Combining search with inference

We can solve a CSP by either:

- perform **pure search** with the backtracking algorithm
- perform **pure inference** (as shown later)

Both methods are correct but do not scale up to big problems

State-of-the-art solvers combine search and **limited but efficient** forms of inference in order to reduce the search space

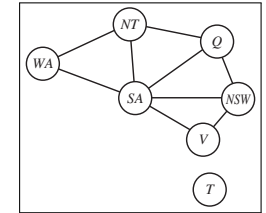
Forward checking

Each node in search tree keeps **current domains** for unassigned variables

Whenever variable X_i is assigned, FC looks at each unassigned variable X_j that is connected to X_i by a constraint, and deletes from D_j all values that are **inconsistent** with value chosen for X_i

Partner of MRV heuristic: select the next variable to assign as one with smallest current domain

Example: Backtracking with forward checking

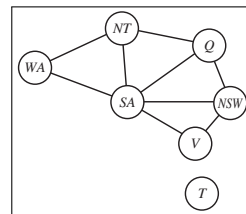


Variable selection with **MRV heuristic**

	WA	NT	Q	NSW	V	SA	T
Initial domains	RGB	RGB	RGB	RGB	RGB	RGB	RGB
After $SA = R$	GB	GB	GB	GB	GB	R	RGB
After $NT = G$	B	G	B	GB	GB	R	RGB
After $Q = B$	B	G	B	G	GB	R	RGB
After $NSW = G$	B	G	B	G	B	R	RGB
After $WA = B$	B	G	B	G	B	R	RGB
After $V = B$	B	G	B	G	B	R	RGB
After $T = R$	B	G	B	G	B	R	R

Example: Backtracking with forward checking

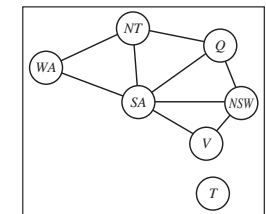
Variable/value selection: $WA = R$, $Q = G$, ...



	WA	NT	Q	NSW	V	SA	T
Initial domains	RGB	RGB	RGB	RGB	RGB	RGB	RGB
After $WA = R$	R	GB	RGB	RGB	RGB	GB	RGB
After $Q = G$	R	B	G	R B	RGB	B	RGB
After $V = B$	R	B	G	R	B	—	RGB
***** BACKTRACK *****							

Example: Backtracking with forward checking


Variable/value selection: $WA = R$, $Q = G$, ...



	WA	NT	Q	NSW	V	SA	T
Initial domains	RGB	RGB	RGB	RGB	RGB	RGB	RGB
After $WA = R$	R	GB	RGB	RGB	RGB	GB	RGB
After $Q = G$	R	B	G	R B	RGB	B	RGB
After $V = R$	R	B	G	B	R	B	RGB
After $NT = B$	R	B	G	B	R	—	RGB
***** BACKTRACK *****							

Example: Backtracking with forward checking


Variable/value selection: $WA = R, Q = G, \dots$



	WA	NT	Q	NSW	V	SA	T
Initial domains	RGB	RGB	RGB	RGB	RGB	RGB	RGB
After $WA = R$	R	GB	RGB	RGB	RGB	GB	RGB
After $Q = G$	R	B	G	R B	RGB	B	RGB
After $V = G$	R	B	G	B	G	B	RGB
After $T = R$	R	B	G	B	G	B	R
After $NT = B$	R	B	G	B	G	—	R
***** BACKTRACK *****							

Example: Backtracking with forward checking


Variable/value selection: $WA = R, Q = G, \dots$



	WA	NT	Q	NSW	V	SA	T
Initial domains	RGB	RGB	RGB	RGB	RGB	RGB	RGB
After $WA = R$	R	GB	RGB	RGB	RGB	GB	RGB
After $Q = G$	R	B	G	R B	RGB	B	RGB
After $V = G$	R	B	G	B	G	B	RGB
After $T = G$	R	B	G	B	G	B	G
After $NT = B$	R	B	G	B	G	—	G
***** BACKTRACK *****							

Example: Backtracking with forward checking

Variable/value selection: $WA = R, Q = G, \dots$



	WA	NT	Q	NSW	V	SA	T
Initial domains	RGB	RGB	RGB	RGB	RGB	RGB	RGB
After $WA = R$	R	GB	RGB	RGB	RGB	GB	RGB
After $Q = G$	R	B	G	R B	RGB	B	RGB
After $V = G$	R	B	G	B	G	B	RGB
After $T = B$	R	B	G	B	G	B	B
After $NT = B$	R	B	G	B	G	—	B
***** BACKTRACK *****							

Chronological and non-chronological backtracking

When search reaches **terminal node** that doesn't correspond to complete assignment (i.e. **conflict node**), the search **backtracks** to **most recent decision point**

Most recent decision point may not be **reason for conflict**

A better idea is to **analyze the conflict** and backtrack to most recent decision point that caused the conflict

Such backtracking is called **non-chronological conflict-based backtracking** and also **conflict-directed backjumping**

Constraint propagation: Arc consistency

Arc consistency is a **property of CSPs**:

- CSP $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ is **arc consistent** iff for each pair of variables X_i and X_j connected in constraint graph, the **arc** (X_i, X_j) is **consistent in P**
- Arc (X_i, X_j) is consistent in P iff for each value x_i of X_i , there exists a value x_j of X_j such that the partial assignment $(X_i = x_i, X_j = x_j)$ is consistent with all constraints (i.e. it doesn't violate any constraint)

For each satisfiable CSP P , there is a CSP P' equivalent to P and with the same variables as P that is arc consistent

An algorithm for arc consistency transforms P into equivalent P' or detects that P has no solution. There are many such algorithms

Arc consistency: AC3

```
1 bool AC3(csp P)
2   Queue Q
3   Insert in Q all arcs (X,Y) in constraint graph
4   while Q is not empty
5     Let (X,Y) := Q.pop()
6     if reduce-arc(X,Y)
7       if Domain[X] == ∅ then return false
8       foreach Z such that (Z,X) is edge in constraint graph
9         Insert arc (Z,X) in Q
10  return true
11
12 bool reduce-arc(variable X, variable Y)
13   removed := false
14   foreach x in Domain[X]
15     found := false
16     foreach y in Domain[Y]
17       if (X=x,Y=y) satisfies all constraints between X and Y
18         found := true
19         break
20   if not found
21     Remove x from Domain[X]
22     removed := true
23   return removed
```

Analysis of AC3

Consider CSP $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ with n variables, and let $d = \max_i |D_i|$:

- Time for reduce-arc(X,Y) is $O(d^2)$ assuming that takes constant time to check whether partial assignment $(X = x, Y = y)$ is consistent with all constraints
- There are $O(n^2)$ initial insertions in the queue
- Arc (Z, X) is re-inserted when a value of X is removed. Since there are $O(d)$ values for X , arc (Z, X) is re-inserted $O(d)$ times
- Number of iterations bounded by $O(n^2 + n^2d) = O(n^2d)$
- **Total time is $O(n^2d^3)$**

Combining search with AC3

Two ways of combining search with AC3:

- Before search starts: make CSP arc-consistent and then do search
- During search: enforce arc consistency at each node during search (known as **Maintaining Arc Consistency or MAC**)

First option is enough in easy problems while the second is necessary for difficult ones

AC4: Keep track of supports

Algorithm for arc consistency that runs in time $O(n^2d^2)$ which is **optimal** since lower bound $\Omega(n^2d^2)$ holds

Idea:

- Keep counters $n(i, x, j)$ for each constraint with scope $\{X_i, X_j\}$ and value $x \in D_i$ that stores **number of values** of X_j that are consistent with $X_i = x$
- Use queue to track values $X = x$ have **lost support**
- Revise counters efficiently

Arc consistency: AC4

```
1  bool AC4(csp P)
2      Queue Q
3
4      % initialization
5      Calculate value of counters n(X,x,Y). If n(X,x,Y) = 0,
6      remove x from Domain[X] and enqueue pair (X,x)
7
8      while Q is not empty
9          Let (X,x) := Q.pop()
10
11         if Domain[X] is empty then
12             return false                                % CSP has no solution
13
14         % value x was removed from Domain[X]
15         foreach (Z,X)
16             foreach z in Domain[Z]
17                 if (Z=z,X=x) is consistent then
18                     Decrement counter n(Z,z,X)
19                     if n(Z,z,X) == 0 then
20                         Remove z from Domain[Z]
21                         Enqueue (Z,z) in Q
22
23     return true
```

Analysis of AC4

Consider CSP $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ with n variables, and let $d = \max_i |D_i|$:

- Time for initialization is $O(n^2d^2)$
- Time of inner loop is $O(nd)$
- Pair (X, x) is added to queue when value x is removed from D_X . Maximum number of pairs in Q is thus $O(nd)$
- **Total time is** $O(n^2d^2 + n^2d^2) = O(n^2d^2)$

Inference for CSPs

We show how to solve CSPs using **pure inference**

Along the way, identify **tractable subclasses** of CSPs that can be solved in polynomial time

High-order consistency

Arc consistency can be generalized to k -consistency

CSP P is **k -consistent** iff for any set of $k - 1$ variables and each consistent assignment for them, the assignment can be **consistently extended** over any other variable

Under this definition:

- P is 1-consistent iff for each variable X and each unary constraint C for X , each value x for X satisfies C
- P is 2-consistent iff P is arc consistent
- ...

P is **strongly k -consistent** iff it is i -consistent for $i = 1, 2, \dots, k$

Examples: Queens

Q			
	Q		

Arc-consistent but not 3-consistent

		Q		
				Q
Q				

2- and 3-consistent but not 4-consistent

Establishing k -consistency (naive algorithm)

```

1  bool k-consistency(csp P)
2      change := true
3      while change
4          change := false
5          foreach subset S of k-1 variables
6              foreach variable X not in S
7                  change := change || k-revise(S,X)
8
9      if domain of some variable is empty then
10         return false
11     else
12         return true
13
14  bool k-revise(S,X)
15      change := false
16      foreach consistent valuation v of S
17          if there is no value x for X such that {v,X=x} is consistent
18              Mark valuation v as forbidden
19              change := true
20      return change
    
```

Remarks on establishing k -consistency

If there is no constraint in which forbidden valuation ν can be eliminated, algorithm is essentially discovering an **implied constraint**

If CSP has only binary constraints, after establishing k -consistency new constraints of order $k - 1$ may appear

Establishing k -consistency takes time $O((2nd)^{2k})$ where n is number of variables and d is maximum cardinality of domains

i -consistency does not imply j -consistency for $j < k$

Solving CSPs by pure inference

Let $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ be a CSP with n variables that is **strongly n -consistent**

The following **backtrack-free** algorithm finds a solution for P or determines P has no solution

1. Let X_1, X_2, \dots, X_n be order for variables (any order will do), and let ν be **empty partial assignment**
2. If domain of X_1 is empty, return **FAILURE**
3. For $i = 1, 2, \dots, n$:
 - **Select value** x_i for X_i that is **consistent** with partial valuation ν
 - **Extend** partial valuation ν with $X_i = x_i$
4. Return valuation ν

Correctness of inference algorithm

We show that a value x_i for X_i that is consistent with the current valuation ν can be found for $i = 1, 2, \dots, n$ (step 4):

- Claim is true for first iteration as ν is the empty valuation, the problem is 1-consistent, and $D_1 \neq \emptyset$
- Consider the $(i + 1)$ th iteration and let ν be current partial valuation at beginning of $(i + 1)$ th iteration. By induction, ν is **consistent**

By strong n -consistency, problem is $(i + 1)$ -consistent. Therefore, any consistent valuation for $\{X_1, \dots, X_i\}$, like ν , can be extended into consistent valuation for any other variable, like X_{i+1}

Then, there is a value x_{i+1} for X_{i+1} that is consistent with ν and the valuation can be extended with $X_{i+1} = x_{i+1}$

At the end ν is a **complete and consistent** assignment; i.e. ν is a solution

Strong consistency and existence of solutions

Let $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ be a CSP with n variables

If P is strongly n -consistent and domain of some variable is non-empty, then P has solution

Tree structure

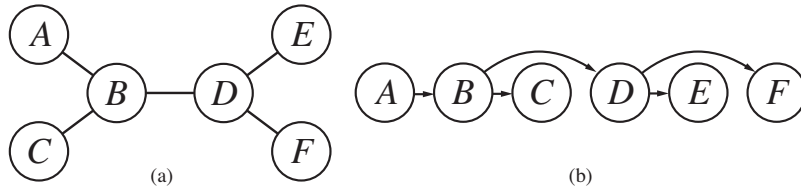
If constraint graph is tree, CSP can be solved in $O(nd^3)$ time

1. Designate any vertex in constraint graph as root and order the vertices (variables) **topologically** so that each vertex appears in the order **after its parent** (it can be done since graph is tree)
2. Enforce **strong arc consistency** in $O(nd^3)$ time (trees have $O(n)$ edges)
3. If domain of first variable is empty, return **FAILURE**
4. Assign values from first to last variable in the order in **backtrack-free** manner as before:

X_1 can be assigned because the problem is 1-consistent and $D_1 \neq \emptyset$

At stage $i + 1$ for X_{i+1} , variable X_{i+1} has only one parent X_j with $j < i$ (as the graph is tree). Since problem is 2-consistent, current assignment can be consistently extended with $X_{i+1} = x_{i+1}$ for some $x_{i+1} \in D_{i+1}$

Topological sort of a tree



[Image from Russell & Norvig. *Artificial Intelligence: A Modern Approach*]

Improved algorithm for tree structure

We can improve algorithm by using **directed arc consistency**

CSP is **directed arc consistent** for order (X_1, X_2, \dots, X_n) iff every arc (X_i, X_j) in constraint graph, for $i < j$, is consistent

1. Topologically order variables as before as (X_1, X_2, \dots, X_n)
2. (Make problem directed arc consistent.) For $j = n$ to 2:
 - Call `reduce-arc(parent($X[j]$), $X[j]$)` to make arc $(\text{parent}(X_j), X_j)$ consistent
 - If domain of $\text{parent}(X_j)$ is empty, return **FAILURE**
3. (Construct valuation in backtrack-free manner.) For $i = 1$ to n :
 - Select value x_i for X_i that is **consistent** with assignment of $\text{parent}(X_i)$. This can be done because X_i has unique parent X_j , with $j < i$, and the directed arc consistency established in step 2

Analysis: each of the $O(n)$ calls to `reduce-arc()` takes time $O(d^2)$. The other steps are done in linear time. **Total time is thus $O(nd^2)$**

Directional consistency

Strong n -consistency is more than needed as variables are assigned values along **fixed variable ordering**

Like improved algorithm for trees, we can enforce appropriate level of consistency along fixed ordering

Width of graph

Let $G = (V, E)$ be undirected graph and \prec be order relation on V :

- \prec -width of vertex v in V : #edges from v to \prec -smaller vertices
- \prec -width of G : maximum \prec -width of vertex in G
- width of G : minimum \prec -width of G over all possible orderings \prec

Width of CSP P is width of its constraint graph

Improved inference algorithm

Let $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ be CSP with constraint graph G

If P is strongly k -consistent, P has width $< k$, and all domains are non-empty, then P has solution

1. Let \prec be ordering on \mathcal{X} such that G has \prec -width $\leq k - 1$
2. Let X_1, X_2, \dots, X_n be \prec -ordering for variables and ν be empty valuation
3. If domain of X_1 is empty, return **FAILURE**
4. For $i = 1, 2, \dots, n$:
 - **Select value** x_i for X_i that is **consistent** with partial valuation ν
 - **Extend** partial valuation ν with $X_i = x_i$
5. Return valuation ν

Remarks for improved inference algorithm

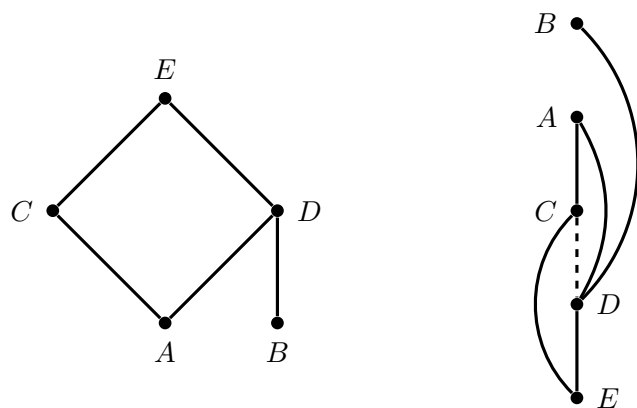
Requires strong k -consistency instead of strong n -consistency ($k \leq n$)

Enforcing strong k -consistency on CSP P may **increase width** of P

We want:

- Selecting variable ordering **dynamically**
- Adjust consistency of each node in **adaptive way**
- Handle width increments in sound manner

Example: Adaptive consistency



Ordering (E, D, C, A, B)

Dechter and Pearl's adaptive consistency

Let $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ be CSP and (X_1, \dots, X_n) be ordering of \mathcal{X}

1. For $i = n, \dots, 1$ do steps (2)–(5)
2. If domain X_i is empty, return **FAILURE**
3. Compute $Parents(X_i) = \{X_j : j < i \text{ and } X_j \text{ is connected to } X_i\}$
4. Add edges between all pairs of variables in $Parents(X_i)$ (among those not already connected)
5. Perform consistency($Parents(X_i), X_i$)
6. Find solution (or determine none exists) in **backtrack-free** manner along order (X_1, \dots, X_n)

Ordering doesn't need to fixed a priori, a good ordering can be **discovered** along execution; obtaining best ordering is **NP-hard**

Other approaches

- “Remove” variables until constraint graph becomes tree that can be solved by algorithm for trees. This is called **cutset conditioning**
- Construct a **tree decomposition** of CSP made of independent subproblems, solve each subproblem independently, and combine solutions into global solution

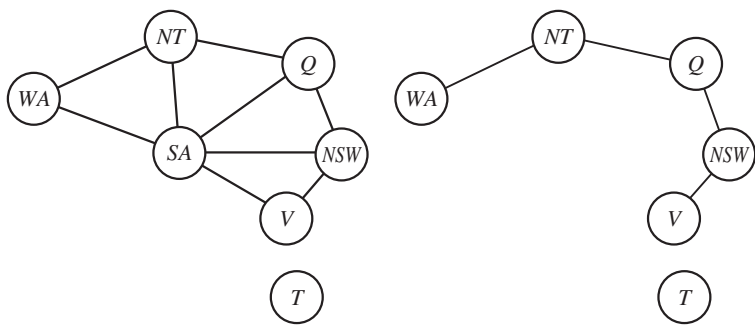
Cutset conditioning

1. **Choose** set S of variables such that after their removal, the constraint graph becomes a tree. S is called **cycle cutset** of constraint graph
2. For each valuation $\nu = \nu_S$ of S , **reduce** P into P_ν by **instantiating** variables in S to values in ν
3. **Solve** P_ν and return overall solution if found
4. If there is no valuation $\nu = \nu_S$ such that P_ν is solvable, return FAILURE
5. If $|S| = c$, reduced CSP can be solved in time $O((n - c)d^2)$ using **directed arc consistency**. Since there are $O(d^c)$ valuations for S , overall algorithm takes time $O((n - c)d^{2+c})$

There is no a priori bound on the size S of a minimum cycle cutset

Finding cycle cutset of **minimum size** is **NP-hard**

Cycle cutset in example




[Image from Russell & Norvig, *Artificial Intelligence: A Modern Approach*]

Tree decomposition

Tree decomposition of CSP $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ is collection of subproblems where each subproblem, defined over subset of variables, is such that:

- Each variable appears in at least one subproblem
- Each constraint $C \in \mathcal{C}$, there is at least one subproblem whose set of variables contains the scope of C
- Subproblems sharing variables are organized into **tree structure**
- If variable X_i appears in two subproblems, X_i then appears in each subproblem along the **unique path** that connects both subproblems

A tree decomposition of the example



© 2018 Blai Bonet

[Image from Russell & Norvig, *Artificial Intelligence: A Modern Approach*]

Analysis

Let $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ be CSP, T be tree decomposition for P with k subproblems, and c be **maximum subproblem size**

- Constructing P' takes time $O(kd^c)$ as there are k subproblems and each subproblem involves $O(d^c)$ valuations over its variables
- Problem P' has k variables, each domain has size $O(d^c)$, and P' has tree structure
- P' can be solved by **directed arc consistency** in time $O(kd^{2c})$
- Total time is thus $O(kd^{2c})$

There is no a priori bound on the maximum subproblem size

Finding **best tree decomposition** is **NP-hard**

© 2018 Blai Bonet

Solving CSPs by tree decompositions

Given CSP $P = (\mathcal{X}, \mathcal{D}, \mathcal{C})$ and tree decomposition for P , construct new **binary CSP** $P' = (\mathcal{X}', \mathcal{D}', \mathcal{C}')$ as follows:

- There is **one variable for each subproblem** in tree decomposition; the i th subproblem corresponds to variable X'_i
- Domain D'_i for variable X'_i corresponds to **all solutions of the i th subproblem** (i th subproblem is viewed as a reduced CSP)
- If i th and j th subproblems are connected (because they share at least one variable), there is **binary constraint** in \mathcal{D}' with scope (X'_i, X'_j) and relation given by all tuples (t'_i, t'_j) such that
 - $t'_i \in D'_i$ and $t'_j \in D'_j$
 - $t'_i[X_k] = t'_j[X_k]$ for every variable X_k that appears in both subproblems (i.e. solutions to subproblems must agree on shared variables)

© 2018 Blai Bonet

Summary

- CSP is a fundamental problem in AI
- CSPs with binary constraints are universal
- CSPs are intractable in general
- CSPs can be solved by either pure search or pure inference
- State-of-the-art solvers combine search with limited and efficient forms of inference

© 2018 Blai Bonet