

Principles of AI Planning

5. Planning as search: progression and regression

Bernhard Nebel and Robert Mattmüller

Albert-Ludwigs-Universität Freiburg

November 4th, 2011

Principles of AI Planning

November 4th, 2011 — 5. Planning as search: progression and regression

5.1 Planning as (classical) search

5.2 Progression

5.3 Regression

5.1 Planning as (classical) search

- Introduction
- Classification of search-based planners

What do we mean by search?

- ▶ **Search** is a very generic term.
- ↪ Every algorithm that tries out various alternatives can be said to “search” in some way.
- ▶ Here, we mean **classical search** algorithms.
 - ▶ **Search nodes** are **expanded** to generate **successor nodes**.
 - ▶ **Examples**: breadth-first search, A*, hill-climbing, ...
- ▶ To be brief, we just say **search** in the following (not “classical search”).

Do you know this stuff already?

- ▶ We **assume prior knowledge** of basic search algorithms:
 - ▶ uninformed vs. informed
 - ▶ systematic vs. local
- ▶ There will be a small refresher in the next chapter.
- ▶ **Background:** Russell & Norvig, Artificial Intelligence – A Modern Approach, Ch. 3 (all of it), Ch. 4 (local search)

Search in planning

- ▶ **search:** one of the **big success stories** of AI
- ▶ many planning algorithms based on classical AI search (we'll see some other algorithms later, though)
- ▶ will be the focus of this and the following chapters (the majority of the course)

Satisficing or optimal planning?

Must carefully distinguish two different problems:

- ▶ **satisficing planning:** any solution is OK (although shorter solutions typically preferred)
- ▶ **optimal planning:** plans must have shortest possible length

Both are often solved by search, but:

- ▶ details are **very different**
- ▶ almost **no overlap** between good techniques for satisficing planning and good techniques for optimal planning
- ▶ many problems that are trivial for satisficing planners are impossibly hard for optimal planners

Planning by search

How to apply search to planning? \rightsquigarrow **many choices to make!**

Choice 1: Search direction

- ▶ **progression:** forward from initial state to goal
- ▶ **regression:** backward from goal states to initial state
- ▶ **bidirectional search**

Planning by search

How to apply search to planning? \rightsquigarrow many choices to make!

Choice 2: Search space representation

- ▶ search nodes are associated with **states**
(\rightsquigarrow **state-space search**)
- ▶ search nodes are associated with **sets of states**

Planning by search

How to apply search to planning? \rightsquigarrow many choices to make!

Choice 3: Search algorithm

- ▶ **uninformed search**:
depth-first, breadth-first, iterative depth-first, ...
- ▶ **heuristic search (systematic)**:
greedy best-first, A*, Weighted A*, IDA*, ...
- ▶ **heuristic search (local)**:
hill-climbing, simulated annealing, beam search, ...

Planning by search

How to apply search to planning? \rightsquigarrow many choices to make!

Choice 4: Search control

- ▶ **heuristics** for informed search algorithms
- ▶ **pruning techniques**: invariants, symmetry elimination, partial-order reduction, helpful actions pruning, ...

Search-based satisficing planners

FF (Hoffmann & Nebel, 2001)

- ▶ **search direction**: forward search
- ▶ **search space representation**: single states
- ▶ **search algorithm**: enforced hill-climbing (informed local)
- ▶ **heuristic**: FF heuristic (inadmissible)
- ▶ **pruning technique**: helpful actions (incomplete)

\rightsquigarrow one of the best satisficing planners

Search-based optimal planners

Fast Downward Stone Soup (Helmert et al., 2011)

- ▶ search direction: forward search
- ▶ search space representation: single states
- ▶ search algorithm: A* (informed systematic)
- ▶ heuristic: multiple admissible heuristics combined into a heuristic portfolio (LM-cut, M&S, blind, ...)
- ▶ pruning technique: none

↪ one of the best optimal planners

Our plan for the next lectures

Choices to make:

1. search direction: progression/regression/both
↪ this chapter
2. search space representation: states/sets of states
↪ this chapter
3. search algorithm: uninformed/heuristic; systematic/local
↪ next chapter
4. search control: heuristics, pruning techniques
↪ following chapters

5.2 Progression

- Overview
- Example

Planning by forward search: progression

Progression: Computing the successor state $app_o(s)$ of a state s with respect to an operator o .

Progression planners find solutions by forward search:

- ▶ start from initial state
- ▶ iteratively pick a previously generated state and **progress it** through an operator, generating a new state
- ▶ solution found when a goal state generated

pro: very easy and efficient to implement

Search space representation in progression planners

Two alternative search spaces for progression planners:

1. search nodes correspond to states

- ▶ when the same state is generated along different paths, it is not considered again (**duplicate detection**)
- ▶ **pro**: save time to consider same state again
- ▶ **con**: memory intensive (must maintain **closed list**)


2. search nodes correspond to operator sequences

- ▶ different operator sequences may lead to identical states (**transpositions**); search does not notice this
- ▶ **pro**: can be very memory-efficient
- ▶ **con**: much wasted work (often exponentially slower)

⇒ first alternative usually preferable in planning
(**unlike** many classical search benchmarks like 15-puzzle)


Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)




Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)




Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)




Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)




Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)




Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)




Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)




Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)




Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)



Progression planning example (depth-first search)

Example where search nodes correspond to operator sequences (no duplicate detection)



5.3 Regression

- Overview
- Example
- Regression for STRIPS tasks
- Regression for general planning tasks
- Practical issues

Forward search vs. backward search

Going through a transition graph in forward and backward directions is **not symmetric**:

- ▶ forward search starts from a **single** initial state;
backward search starts from a **set** of goal states
- ▶ when applying an operator o in a state s in forward direction, there is a **unique successor state** s' ;
if we applied operator o to end up in state s' ,
there can be **several possible predecessor states** s

↪ most natural representation for backward search in planning associates **sets of states** with search nodes

Planning by backward search: regression

Regression: Computing the possible predecessor states $regr_o(G)$ of a set of states G with respect to the last operator o that was applied.

Regression planners find solutions by backward search:

- ▶ start from set of goal states
- ▶ iteratively pick a previously generated state set and **regress it** through an operator, generating a new state set
- ▶ solution found when a generated state set includes the initial state

Pro: can handle many states simultaneously


Con: basic operations complicated and expensive

Search space representation in regression planners

identify state sets with **logical formulae** (again):

- ▶ **search nodes correspond to state sets**
- ▶ each state set is represented by a **logical formula**:
 φ represents $\{s \in S \mid s \models \varphi\}$
- ▶ many basic search operations like detecting duplicates are NP-hard or coNP-hard

Regression planning example (depth-first search)




Regression planning example (depth-first search)

 γ 

Regression planning example (depth-first search)

$$\varphi_1 = \text{regr}_{\rightarrow}(\gamma)$$

$$\varphi_1 \rightarrow \gamma$$




Regression planning example (depth-first search)

$$\varphi_1 = \text{regr}_{\rightarrow}(\gamma)$$

$$\varphi_2 = \text{regr}_{\rightarrow}(\varphi_1)$$

$$\varphi_2 \rightarrow \varphi_1 \rightarrow \gamma$$




Regression planning example (depth-first search)

$$\varphi_1 = \text{regr}_{\rightarrow}(\gamma)$$

$$\varphi_2 = \text{regr}_{\rightarrow}(\varphi_1)$$

$$\varphi_3 = \text{regr}_{\rightarrow}(\varphi_2), I \models \varphi_3$$

$$\varphi_3 \rightarrow \varphi_2 \rightarrow \varphi_1 \rightarrow \gamma$$



Regression for STRIPS planning tasks

Definition (STRIPS planning task)

A planning task is a **STRIPS planning task** if all operators are STRIPS operators and the goal is a conjunction of atoms.

Regression for **STRIPS planning tasks** is very simple:

- ▶ Goals are conjunctions of atoms $a_1 \wedge \dots \wedge a_n$.
- ▶ **First step:** Choose an operator that makes none of a_1, \dots, a_n false.
- ▶ **Second step:** Remove goal atoms achieved by the operator (if any) and add its preconditions.

↪ Outcome of regression is again conjunction of atoms.

Optimization: only consider operators making some a_i true

STRIPS regression

Definition (STRIPS regression)


Let $\varphi = \varphi_1 \wedge \dots \wedge \varphi_n$ be a conjunction of atoms, and let $o = \langle \chi, e \rangle$ be a STRIPS operator which adds the atoms a_1, \dots, a_k and deletes the atoms d_1, \dots, d_l .

The **STRIPS regression** of φ with respect to o is

$$\text{sregr}_o(\varphi) := \begin{cases} \perp & \text{if } a_i = d_j \text{ for some } i, j \\ \perp & \text{if } \varphi_i = d_j \text{ for some } i, j \\ \chi \wedge \bigwedge (\{\varphi_1, \dots, \varphi_n\} \setminus \{a_1, \dots, a_k\}) & \text{otherwise} \end{cases}$$

Note: $\text{sregr}_o(\varphi)$ is again a conjunction of atoms, or \perp .

STRIPS regression example



Note: Predecessor states are in general not unique.
This picture is just for illustration purposes.

$$\begin{aligned} o_1 &= \langle \text{on} \text{red} \wedge \text{on} \text{blue}, \neg \text{on} \text{red} \wedge \neg \text{on} \text{blue} \wedge \text{on} \text{green} \rangle \\ o_2 &= \langle \text{on} \text{red} \wedge \text{on} \text{green} \wedge \text{on} \text{blue}, \neg \text{on} \text{red} \wedge \neg \text{on} \text{green} \wedge \neg \text{on} \text{blue} \rangle \\ o_3 &= \langle \text{on} \text{red} \wedge \text{on} \text{green}, \neg \text{on} \text{red} \wedge \neg \text{on} \text{green} \wedge \text{on} \text{blue} \rangle \end{aligned}$$

$$\gamma = \text{on} \text{red} \wedge \text{on} \text{green}$$

$$\varphi_1 = \text{sregr}_{o_3}(\gamma) = \text{on} \text{red} \wedge \text{on} \text{green} \wedge \text{on} \text{blue}$$

$$\varphi_2 = \text{sregr}_{o_2}(\varphi_1) = \text{on} \text{red} \wedge \text{on} \text{green} \wedge \text{on} \text{blue} \wedge \text{on} \text{red}$$

$$\varphi_3 = \text{sregr}_{o_1}(\varphi_2) = \text{on} \text{red} \wedge \text{on} \text{green} \wedge \text{on} \text{blue} \wedge \text{on} \text{red} \wedge \text{on} \text{red}$$

Regression for general planning tasks

- ▶ With disjunctions and conditional effects, things become more tricky.
How to regress $a \vee (b \wedge c)$ with respect to $\langle q, d \triangleright b \rangle$?
- ▶ The story about goals and subgoals and fulfilling subgoals, as in the STRIPS case, is no longer useful.
- ▶ We present a general method for doing regression for any formula and any operator.
- ▶ Now we extensively use the idea of representing sets of states as formulae.

Effect preconditions

Definition (effect precondition)

The **effect precondition** $EPC_I(e)$ for literal I and effect e is defined as follows:

$$\begin{aligned} EPC_I(I) &= \top \\ EPC_I(I') &= \perp \text{ if } I \neq I' \text{ (for literals } I') \\ EPC_I(e_1 \wedge \dots \wedge e_n) &= EPC_I(e_1) \vee \dots \vee EPC_I(e_n) \\ EPC_I(\chi \triangleright e) &= EPC_I(e) \wedge \chi \end{aligned}$$

Intuition: $EPC_I(e)$ describes the situations in which effect e causes literal I to become true.

Effect precondition examples

Example

$$\begin{aligned} EPC_a(b \wedge c) &= \perp \vee \perp \equiv \perp \\ EPC_a(a \wedge (b \triangleright a)) &= \top \vee (\top \wedge b) \equiv \top \\ EPC_a((c \triangleright a) \wedge (b \triangleright a)) &= (\top \wedge c) \vee (\top \wedge b) \equiv c \vee b \end{aligned}$$

Effect preconditions: connection to change sets

Lemma (A)

Let s be a state, I a literal and e an effect.

Then $I \in [e]_s$ if and only if $s \models EPC_I(e)$.

Proof.

Induction on the structure of the effect e .

Base case 1, $e = I$: $I \in [I]_s = \{I\}$ by definition, and $s \models EPC_I(I) = \top$ by definition. Both sides of the equivalence are true.

Base case 2, $e = I'$ for some literal $I' \neq I$: $I \notin [I']_s = \{I'\}$ by definition, and $s \not\models EPC_I(I') = \perp$ by definition. Both sides are false.

Effect preconditions: connection to change sets

Proof (ctd.)

Inductive case 1, $e = e_1 \wedge \dots \wedge e_n$:

$$I \in [e]_s \text{ iff } I \in [e_1]_s \cup \dots \cup [e_n]_s \quad (\text{Def } [e_1 \wedge \dots \wedge e_n]_s)$$

$$\text{iff } I \in [e']_s \text{ for some } e' \in \{e_1, \dots, e_n\}$$

$$\text{iff } s \models EPC_I(e') \text{ for some } e' \in \{e_1, \dots, e_n\} \quad (\text{IH})$$

$$\text{iff } s \models EPC_I(e_1) \vee \dots \vee EPC_I(e_n)$$

$$\text{iff } s \models EPC_I(e_1 \wedge \dots \wedge e_n). \quad (\text{Def } EPC)$$

Inductive case 2, $e = \chi \triangleright e'$:

$$I \in [\chi \triangleright e']_s \text{ iff } I \in [e']_s \text{ and } s \models \chi \quad (\text{Def } [\chi \triangleright e']_s)$$

$$\text{iff } s \models EPC_I(e') \text{ and } s \models \chi \quad (\text{IH})$$

$$\text{iff } s \models EPC_I(e') \wedge \chi$$

$$\text{iff } s \models EPC_I(\chi \triangleright e'). \quad (\text{Def } EPC)$$

□

Effect preconditions: connection to normal form

Remark: *EPC* vs. effect normal form

Notice that in terms of $EPC_a(e)$, any operator $\langle \chi, e \rangle$ can be expressed in effect normal form as

$$\left\langle \chi, \bigwedge_{a \in A} ((EPC_a(e) \triangleright a) \wedge (EPC_{\neg a}(e) \triangleright \neg a)) \right\rangle,$$

where A is the set of all state variables.

Regressing state variables

The formula $EPC_a(e) \vee (a \wedge \neg EPC_{\neg a}(e))$ expresses the **value of state variable $a \in A$ after applying o** in terms of **values of state variables before applying o** .

Either:

- ▶ a became true, or
- ▶ a was true before and it did not become false.

Regressing state variables: examples

Example

Let $e = (b \triangleright a) \wedge (c \triangleright \neg a) \wedge b \wedge \neg d$.

variable x	$EPC_x(e) \vee (x \wedge \neg EPC_{\neg x}(e))$
a	$b \vee (a \wedge \neg c)$
b	$\top \vee (b \wedge \neg \perp) \equiv \top$
c	$\perp \vee (c \wedge \neg \perp) \equiv c$
d	$\perp \vee (d \wedge \neg \top) \equiv \perp$

Regressing state variables: correctness

Lemma (B)

Let a be a state variable, $o = \langle \chi, e \rangle$ an operator, s a state, and $s' = \text{app}_o(s)$.

Then $s \models EPC_a(e) \vee (a \wedge \neg EPC_{\neg a}(e))$ if and only if $s' \models a$.

Proof.

(\Rightarrow): Assume $s \models EPC_a(e) \vee (a \wedge \neg EPC_{\neg a}(e))$.

Do a case analysis on the two disjuncts.

1. Assume that $s \models EPC_a(e)$. By Lemma A, we have $a \in [e]_s$ and hence $s' \models a$.
2. Assume that $s \models a \wedge \neg EPC_{\neg a}(e)$. By Lemma A, we have $\neg a \notin [e]_s$. Hence a remains true in s' .

Regressing state variables: correctness

Proof (ctd.)

(\Leftarrow): We showed that if the formula is **true** in s , then a is **true** in s' . For the second part, we show that if the formula is **false** in s , then a is **false** in s' .

- ▶ So assume $s \not\models EPC_a(e) \vee (a \wedge \neg EPC_{\neg a}(e))$.
- ▶ Then $s \models \neg EPC_a(e) \wedge (\neg a \vee EPC_{\neg a}(e))$ (de Morgan).
- ▶ Case distinction: a is true or a is false in s .
 1. Assume that $s \models a$. Now $s \models EPC_{\neg a}(e)$ because $s \models \neg a \vee EPC_{\neg a}(e)$. Hence by Lemma A $\neg a \in [e]_s$ and we get $s' \not\models a$.
 2. Assume that $s \not\models a$. Because $s \models \neg EPC_a(e)$, by Lemma A we get $a \notin [e]_s$ and hence $s' \not\models a$.

Therefore in both cases $s' \not\models a$.

□

Regression: general definition

We base the definition of regression on formulae $EPC_I(e)$.

Definition (general regression)

Let φ be a propositional formula and $o = \langle \chi, e \rangle$ an operator.

The **regression of φ with respect to o** is

$$\text{regr}_o(\varphi) = \chi \wedge \varphi_r \wedge \kappa$$

where

1. φ_r is obtained from φ by replacing each $a \in A$ by $EPC_a(e) \vee (a \wedge \neg EPC_{\neg a}(e))$, and
2. $\kappa = \bigwedge_{a \in A} \neg(EPC_a(e) \wedge EPC_{\neg a}(e))$.

The formula κ expresses that operators are only applicable in states where their change sets are consistent.

Regression examples

- ▶ $\text{regr}_{\langle a, b \rangle}(b) \equiv a \wedge (\top \vee (b \wedge \neg \perp)) \wedge \top \equiv a$
- ▶ $\text{regr}_{\langle a, b \rangle}(b \wedge c \wedge d)$
 $\equiv a \wedge (\top \vee (b \wedge \neg \perp)) \wedge (\perp \vee (c \wedge \neg \perp)) \wedge (\perp \vee (d \wedge \neg \perp)) \wedge \top$
 $\equiv a \wedge c \wedge d$
- ▶ $\text{regr}_{\langle a, c \triangleright b \rangle}(b) \equiv a \wedge (c \vee (b \wedge \neg \perp)) \wedge \top \equiv a \wedge (c \vee b)$
- ▶ $\text{regr}_{\langle a, (c \triangleright b) \wedge (b \triangleright \neg b) \rangle}(b) \equiv a \wedge (c \vee (b \wedge \neg b)) \wedge \neg(c \wedge b)$
 $\equiv a \wedge c \wedge \neg b$
- ▶ $\text{regr}_{\langle a, (c \triangleright b) \wedge (d \triangleright \neg b) \rangle}(b) \equiv a \wedge (c \vee (b \wedge \neg d)) \wedge \neg(c \wedge d)$
 $\equiv a \wedge (c \vee b) \wedge (c \vee \neg d) \wedge (\neg c \vee \neg d)$
 $\equiv a \wedge (c \vee b) \wedge \neg d$

Regression example: binary counter

$$\begin{aligned} &(\neg b_0 \triangleright b_0) \wedge \\ &((\neg b_1 \wedge b_0) \triangleright (b_1 \wedge \neg b_0)) \wedge \\ &((\neg b_2 \wedge b_1 \wedge b_0) \triangleright (b_2 \wedge \neg b_1 \wedge \neg b_0)) \end{aligned}$$

$$EPC_{b_2}(e) = \neg b_2 \wedge b_1 \wedge b_0$$

$$EPC_{b_1}(e) = \neg b_1 \wedge b_0$$

$$EPC_{b_0}(e) = \neg b_0$$

$$EPC_{\neg b_2}(e) = \perp$$

$$EPC_{\neg b_1}(e) = \neg b_2 \wedge b_1 \wedge b_0$$

$$EPC_{\neg b_0}(e) = (\neg b_1 \wedge b_0) \vee (\neg b_2 \wedge b_1 \wedge b_0) \equiv (\neg b_1 \vee \neg b_2) \wedge b_0$$

Regression replaces state variables as follows:

$$b_2 \text{ by } (\neg b_2 \wedge b_1 \wedge b_0) \vee (b_2 \wedge \neg \perp) \equiv (\neg b_1 \wedge b_0) \vee b_2$$

$$b_1 \text{ by } (\neg b_1 \wedge b_0) \vee (b_1 \wedge \neg(\neg b_2 \wedge b_1 \wedge b_0))$$

$$\equiv (\neg b_1 \wedge b_0) \vee (b_1 \wedge (b_2 \vee \neg b_0))$$

$$b_0 \text{ by } \neg b_0 \vee (b_0 \wedge \neg((\neg b_1 \vee \neg b_2) \wedge b_0)) \equiv \neg b_0 \vee (b_1 \wedge b_2)$$

General regression: correctness

Theorem (correctness of $\text{regr}_o(\varphi)$)

Let φ be a formula, o an operator and s a state.

Then $s \models \text{regr}_o(\varphi)$ iff o is applicable in s and $\text{app}_o(s) \models \varphi$.

Proof.

Let $o = \langle \chi, e \rangle$. Recall that $\text{regr}_o(\varphi) = \chi \wedge \varphi_r \wedge \kappa$, where φ_r and κ are as defined previously.

If o is inapplicable in s , then $s \not\models \chi \wedge \kappa$, both sides of the “iff” condition are false, and we are done. Hence, we only further consider states s where o is applicable. Let $s' := \text{app}_o(s)$.

We know that $s \models \chi \wedge \kappa$ (because o is applicable), so the “iff” condition we need to prove simplifies to:

$$s \models \varphi_r \text{ iff } s' \models \varphi.$$

General regression: correctness

Proof (ctd.)

To show: $s \models \varphi_r$ iff $s' \models \varphi$.

We show that for all formulae ψ , $s \models \psi_r$ iff $s' \models \psi$, where ψ_r is ψ with every $a \in A$ replaced by $\text{EPC}_a(e) \vee (a \wedge \neg \text{EPC}_{\neg a}(e))$.

The proof is by structural induction on ψ .

Induction hypothesis $s \models \psi_r$ if and only if $s' \models \psi$.

Base cases 1 & 2 $\psi = \top$ or $\psi = \perp$: trivial, as $\psi_r = \psi$.

Base case 3 $\psi = a$ for some $a \in A$:

Then $\psi_r = \text{EPC}_a(e) \vee (a \wedge \neg \text{EPC}_{\neg a}(e))$.

By Lemma B, $s \models \psi_r$ iff $s' \models \psi$.

General regression: correctness

Proof (ctd.)

Inductive case 1 $\psi = \neg\psi'$:

$$\begin{aligned} s \models \psi_r \text{ iff } s \models (\neg\psi')_r \text{ iff } s \models \neg(\psi'_r) \text{ iff } s \not\models \psi'_r \\ \text{iff (IH)} s' \not\models \psi' \text{ iff } s' \models \neg\psi' \text{ iff } s' \models \psi \end{aligned}$$

Inductive case 2 $\psi = \psi' \vee \psi''$:

$$\begin{aligned} s \models \psi_r \text{ iff } s \models (\psi' \vee \psi'')_r \text{ iff } s \models \psi'_r \vee \psi''_r \\ \text{iff } s \models \psi'_r \text{ or } s \models \psi''_r \\ \text{iff (IH, twice)} s' \models \psi' \text{ or } s' \models \psi'' \\ \text{iff } s' \models \psi' \vee \psi'' \text{ iff } s' \models \psi \end{aligned}$$

Inductive case 3 $\psi = \psi' \wedge \psi''$: Very similar to inductive case 2, just with \wedge instead of \vee and “and” instead of “or”.

□

Emptiness and subsumption testing

The following two tests are useful when performing regression searches to avoid exploring unpromising branches:

- ▶ Test that $\text{regr}_o(\varphi)$ does not represent the empty set (which would mean that search is in a dead end).
For example, $\text{regr}_{\langle a, \neg p \rangle}(p) \equiv a \wedge \perp \equiv \perp$.
- ▶ Test that $\text{regr}_o(\varphi)$ does not represent a subset of φ (which would make the problem harder than before).
For example, $\text{regr}_{\langle b, c \rangle}(a) \equiv a \wedge b$.

Both of these problems are **NP-hard**.

Formula growth

The formula $\text{regr}_{o_1}(\text{regr}_{o_2}(\dots \text{regr}_{o_{n-1}}(\text{regr}_{o_n}(\varphi))))$ may have size $O(|\varphi||o_1||o_2|\dots|o_{n-1}||o_n|)$, i.e., the product of the sizes of φ and the operators.

\rightsquigarrow worst-case **exponential** size $O(m^n)$

Logical simplifications

- ▶ $\perp \wedge \varphi \equiv \perp$, $\top \wedge \varphi \equiv \varphi$, $\perp \vee \varphi \equiv \varphi$, $\top \vee \varphi \equiv \top$
- ▶ $a \vee \varphi \equiv a \vee \varphi[\perp/a]$, $\neg a \vee \varphi \equiv \neg a \vee \varphi[\top/a]$, $a \wedge \varphi \equiv a \wedge \varphi[\top/a]$, $\neg a \wedge \varphi \equiv \neg a \wedge \varphi[\perp/a]$
- ▶ idempotency, absorption, commutativity, associativity, ...

Restricting formula growth in search trees

Problem very big formulae obtained by regression


Cause **disjunctivity** in the (NNF) formulae
(formulae **without disjunctions** easily convertible to small formulae $l_1 \wedge \dots \wedge l_n$ where l_i are literals and n is at most the number of state variables.)

Idea handle disjunctivity when generating search trees

Unrestricted regression: search tree example

Unrestricted regression: do not treat disjunctions specially

Goal $\gamma = a \wedge b$, initial state $I = \{a \mapsto 0, b \mapsto 0, c \mapsto 0\}$.




Full splitting: search tree example

Full splitting: always remove all disjunctivity

Goal $\gamma = a \wedge b$, initial state $I = \{a \mapsto 0, b \mapsto 0, c \mapsto 0\}$.

$(\neg c \vee a) \wedge b$ in DNF: $(\neg c \wedge b) \vee (a \wedge b)$

\rightsquigarrow split into $\neg c \wedge b$ and $a \wedge b$



General splitting strategies

Alternatives:

1. Do nothing (**unrestricted regression**).
2. Always eliminate all disjunctivity (**full splitting**).
3. Reduce disjunctivity if formula becomes too big.

Discussion:

- ▶ With **unrestricted regression** the formulae may have **size that is exponential** in the number of state variables.
- ▶ With **full splitting** search tree can be **exponentially bigger** than without splitting.
- ▶ The third option lies between these two extremes.

Summary

- ▶ (Classical) **search** is a very important planning approach.
- ▶ Search-based planning algorithms differ along many dimensions, including
 - ▶ **search direction** (forward, backward)
 - ▶ **what each search node represents** (a state, a set of states, an operator sequence)
- ▶ **Progression search** proceeds forwards from the initial state.
 - ▶ If we use duplicate detection, each search node corresponds to a unique **state**.
 - ▶ If we do not use duplicate detection, each search node corresponds to a unique **operator sequence**.

Summary (ctd.)

- ▶ **Regression search** proceeds backwards from the goal.
 - ▶ Each search node corresponds to a **set of states** represented by a **formula**.
 - ▶ Regression is simple for **STRIPS** operators.
 - ▶ The theory for **general regression** is more complex.
 - ▶ When applying regression in practice, additional considerations such as when and how to perform **splitting** come into play.