

Learning Features and Abstract Actions for Computing Generalized Plans

Blai Bonet¹ Guillem Francès² Hector Geffner³

¹Universidad Simón Bolívar, Caracas, Venezuela

²University of Basel, Basel, Switzerland

³ICREA & Universitat Pompeu Fabra, Barcelona, Spain



University
of Basel



ICREA



Universitat
Pompeu Fabra
Barcelona

Planning and Generalized Planning

- Planning is about solving **single planning instances**
- **Generalized planning** is about solving **multiple instances** at once

For example:

- Achieve goal $on(x, y)$ in Blocksworld (any number of blocks, any configuration)
- Go to target location (x^*, y^*) in empty square grid of any size
- Pick objects in grid (any number and location, any grid size)

Srivastava et al. 2008, B. et al. 2009, Hu & De Giacomo 2011, Illanes & McIlraith 2017, ...

Example: Plan for $clear(x)$ using Right Abstraction

- Get $clear(x)$, for designated block x , on **any Blockworld instance**
- **Features:** $F = \{H, n\}$ where
 - H is **Boolean feature** that tells whether gripper is holding a block
 - n is **numerical feature** that counts number of blocks above x
- **Abstract actions:** $A_F = \{\text{Pick-above-}x, \text{Put-aside}\}$ given by
 - $\text{Pick-above-}x = \langle \neg H, n > 0; H, n \downarrow \rangle$
 - $\text{Put-aside} = \langle H; \neg H \rangle$
- **Solution:** If $\neg H \wedge n > 0$ then $\text{Pick-above-}x$; If H then Put-aside
- **Computed** with off-the-shelf FOND planner

Can we learn the RIGHT abstraction?

Features for Generalized Planning: Requirements

Features required for solving collection \mathcal{Q} of instances:

- Must be **general**; i.e. well defined on any state for any instance in \mathcal{Q}
- Must be **predictable**; i.e. effects of actions on features is predictable
- Must **distinguish** goal from non-goal states

Solving all instances in \mathcal{Q} is mapped to solving **single FOND problem** over the features

FOND = Fully Observable Non-Deterministic

Learning Features and Abstract Actions using SAT

- **Input:**

- \mathcal{S} = sample of **state transitions** (s, s') for some instances in \mathcal{Q}
- \mathcal{F} = **pool of features** computed from primitive predicates in \mathcal{Q}

Learning Features and Abstract Actions using SAT

- **Input:**

- \mathcal{S} = sample of **state transitions** (s, s') for some instances in \mathcal{Q}
- \mathcal{F} = **pool of features** computed from primitive predicates in \mathcal{Q}

- **Propositional variables:**

- $selected(f)$ for each feature f (f in “final set” F iff $selected(f)$ is true)
- $D_1(s, t)$ iff selected features **distinguish states** s and t in sample \mathcal{S}
- $D_2(s, s', t, t')$ iff selected features **distinguish transitions** (s, s') and (t, t') in \mathcal{S}

Learning Features and Abstract Actions using SAT

- **Input:**

- \mathcal{S} = sample of **state transitions** (s, s') for some instances in \mathcal{Q}
- \mathcal{F} = **pool of features** computed from primitive predicates in \mathcal{Q}

- **Propositional variables:**

- $selected(f)$ for each feature f (f in “final set” F iff $selected(f)$ is true)
- $D_1(s, t)$ iff selected features **distinguish states** s and t in sample \mathcal{S}
- $D_2(s, s', t, t')$ iff selected features **distinguish transitions** (s, s') and (t, t') in \mathcal{S}

- **Formulas:**

- $D_1(s, t) \iff \bigvee_f selected(f)$ (for f that make s and t different)
- $D_2(s, s', t, t') \iff \bigvee_f selected(f)$ (for f that make (s, s') and (t, t') different)

Learning Features and Abstract Actions using SAT

- **Input:**

- \mathcal{S} = sample of **state transitions** (s, s') for some instances in \mathcal{Q}
- \mathcal{F} = **pool of features** computed from primitive predicates in \mathcal{Q}

- **Propositional variables:**

- $selected(f)$ for each feature f (f in “final set” F iff $selected(f)$ is true)
- $D_1(s, t)$ iff selected features **distinguish states** s and t in sample \mathcal{S}
- $D_2(s, s', t, t')$ iff selected features **distinguish transitions** (s, s') and (t, t') in \mathcal{S}

- **Formulas:**

- $D_1(s, t) \iff \bigvee_f selected(f)$ (for f that make s and t different)
- $D_2(s, s', t, t') \iff \bigvee_f selected(f)$ (for f that make (s, s') and (t, t') different)
- $\neg D_1(s, t)$ (for s and t such **only one** is goal)

Learning Features and Abstract Actions using SAT

- **Input:**

- \mathcal{S} = sample of **state transitions** (s, s') for some instances in \mathcal{Q}
- \mathcal{F} = **pool of features** computed from primitive predicates in \mathcal{Q}

- **Propositional variables:**

- $selected(f)$ for each feature f (f in “final set” F iff $selected(f)$ is true)
- $D_1(s, t)$ iff selected features **distinguish states** s and t in sample \mathcal{S}
- $D_2(s, s', t, t')$ iff selected features **distinguish transitions** (s, s') and (t, t') in \mathcal{S}

- **Formulas:**

- $D_1(s, t) \iff \bigvee_f selected(f)$ (for f that make s and t different)
- $D_2(s, s', t, t') \iff \bigvee_f selected(f)$ (for f that make (s, s') and (t, t') different)
- $\neg D_1(s, t)$ (for s and t such **only one** is goal)
- $\bigwedge_{t'} D_2(s, s', t, t') \implies D_1(s, t)$ (for each (s, s') and t in \mathcal{S})

Learning Features and Abstract Actions using SAT

- **Input:**

- \mathcal{S} = sample of **state transitions** (s, s') for some instances in \mathcal{Q}
- \mathcal{F} = **pool of features** computed from primitive predicates in \mathcal{Q}

- **Propositional variables:**

- $selected(f)$ for each feature f (f in “final set” F iff $selected(f)$ is true)
- $D_1(s, t)$ iff selected features **distinguish states** s and t in sample \mathcal{S}
- $D_2(s, s', t, t')$ iff selected features **distinguish transitions** (s, s') and (t, t') in \mathcal{S}

- **Formulas:**

- $D_1(s, t) \iff \bigvee_f selected(f)$ (for f that make s and t different)
- $D_2(s, s', t, t') \iff \bigvee_f selected(f)$ (for f that make (s, s') and (t, t') different)
- $\neg D_1(s, t)$ (for s and t such **only one** is goal)
- $\bigwedge_{t'} D_2(s, s', t, t') \implies D_1(s, t)$ (for each (s, s') and t in \mathcal{S})

- **Guarantee:** Theory $T(\mathcal{S}, \mathcal{F})$ is SAT iff there is **sound and complete** abstraction relative to sample \mathcal{S} (abstraction is easily obtained from model)

Pool of Features

- Pool \mathcal{F} obtained from primitive and newly defined predicates in \mathcal{Q}
- Numerical and Boolean features n and f defined from **unary predicates** $q(\cdot)$:
 - $n(s) = |\{x : s \models q(x)\}|$ (cardinality of set)
 - $f(s) = |\{x : s \models q(x)\}| > 0$ (whether set is empty or not)
- New unary predicates obtained with **concept grammar**
- “Distance notion” also defined with concept grammar using binary predicates
- Feature f has $cost(f)$ given by its “concept complexity”
- MaxSAT solver minimizes $\sum_{f: selected(f)} cost(f)$

We look for most economical abstraction!

Computational Workflow

For solving generalized problem \mathcal{Q} :

1. Sample set of transition \mathcal{S} from some instances in \mathcal{Q}
2. Compute pool of features \mathcal{F} from primitive predicates, grammar, and bounds
3. **MaxSAT** to find model of theory for $(\mathcal{S}, \mathcal{F})$ of min cost $\sum_{f: \text{selected}(f)} \text{cost}(f)$
4. Decode SAT model to extract abstraction
5. Solve abstraction with off-the-shelf FOND planner

Experimental Result: $\mathcal{Q}_{gripper}$

- **Training:** 2 instances with 4 and 5 balls respectively
- **Features learned** ($|\mathcal{S}| = 403, |\mathcal{F}| = 130$):
 - X = “whether robbly is in target room”
 - B = “number of balls not in target room”
 - C = “number of balls carried by robbly”
 - G = “number of empty grippers (available capacity)”
- **Abstract actions learned:**
 - Drop-ball-at-target = $\langle C > 0, X; C \downarrow, G \uparrow \rangle$
 - Move-to-target-fully-loaded = $\langle \neg X, C > 0, G = 0; X \rangle$
 - Move-to-target-half-loaded = $\langle \neg X, B = 0, C > 0, G > 0; X \rangle$
 - Pick-ball-not-in-target = $\langle \neg X, B > 0, G > 0; B \downarrow, G \downarrow, C \uparrow \rangle$
 - Leave-target = $\langle X, C = 0, G > 0; \neg X \rangle$
- FOND solved in 171.92 secs; MaxSAT time is 0.01 secs
- Plan solves instances for **any number of grippers, any number of balls**

Experimental Result: Q_{reward}

- Pick rewards in grid with **blocked cells** (from *Towards Deep Symbolic RL*, Garnelo, Arulkumaran, Shanahan, 2016)
- STRIPS instances with predicates: $reward(\cdot)$, $at(\cdot)$, $blocked(\cdot)$, $adj(\cdot, \cdot)$
- **Training:** 2 instances 4×4 , 5×5 , diff. dist. of blocked cells and rewards
- **Features learned** ($|\mathcal{S}| = 568$, $|\mathcal{F}| = 280$):
 - R = “number of remaining rewards”
 - D = “min distance to closest reward along unblocked path”
- **Abstract actions learned:**
 - Collect = $\langle D = 0, R > 0; R \downarrow, D \uparrow \rangle$
 - Move-to-closest-reward = $\langle R > 0, D > 0; D \downarrow \rangle$
- FOND solved in 1.36 secs; MaxSAT time is 0.01 secs
- Plan solves **any grid size, any number of rewards, any dist. of blocked cells**

Summary and Future

- **Inductive framework** for generalized planning that mixes **learning** and **planning**
 - Learner needs to learn abstraction (not plans)
 - Planner uses abstraction, transformed, to compute general plans
- **Number of samples** is small as learner only identifies features to be tracked
- Unlike purely learning approaches:
 - Features and policies are **transparent**
 - **Scope and correctness** of plans can be formally characterized
- Relation to **dimensionality reduction** and **embeddings** in ML/Deep Learning

FOND translator: <https://github.com/bonetblai/qnp2fond>