

Learning More Expressive General Policies for Classical Planning Domains

Simon Ståhlberg¹ Blai Bonet² Hector Geffner¹

¹RWTH Aachen University, Germany

²Universitat Pompeu Fabra, Spain



Introduction

- General policies represent strategies for solving many planning instances
 - ▷ E.g., general policy for solving **all** Blocksworld problems
- Three main methods for learning such policies (no “synthesis” methods yet!)
 - ▷ **Combinatorial optimization using explicit pool of \mathcal{C}_2 features** obtained from domain predicates [B. *et al.*, 2019; Francès *et al.*, 2021]
 - Transparent, can be proved correct, trouble scaling up
 - ▷ **Deep learning (DL) using domain predicates but no explicit pool** [Toyer *et al.*, 2020; Garg *et al.*, 2020]
 - Opaque, complex, but scalable
 - ▷ **DL exploiting relation between \mathcal{C}_2 logic and GNNs** [Barceló *et al.*, 2020; Grohe, 2020; Ståhlberg *et al.*, 2022]
 - **R-GNN architecture** adapted from Max-CSP[Γ] [Toenshoff *et al.*, 2021]
 - More transparent and simple, scalable
 - Supervised and non-supervised training
 - **Problem:** insufficient expressivity for generalized planning

In this Work

- Novel relational architecture R-GNN[t], with parameter $t \geq 0$, that combines the R-GNN architecture with a parameterized encoding $A_t(S)$ of planning states S
- As t increases, the expressive power of R-GNN[t] increases, approaching the full expressivity of \mathcal{C}_3 logic
- Significant improvements obtained even with $t = 1$, as shown in experiments
- 2- or 3-GNNs and Edge Transformers unfeasible in practice and limited to binary relations:
 - ▷ 2-GNNs: $\Theta(N^2)$ memory, $\Theta(N^3)$ time, \mathcal{C}_2 expressivity (yet see below)
 - ▷ 3-GNNs: $\Theta(N^3)$ memory, $\Theta(N^4)$ time, \mathcal{C}_3 expressivity
 - ▷ ETs: $\Theta(N^2)$ memory, $\Theta(N^3)$ time, \mathcal{C}_3 expressivity [Müller *et al.*, 2024]

Generalized Planning and First-Order STRIPS

- Generalized planning is about finding **general policies** that solve classes of planning problems
- Task is collection $\{P_1, P_2, P_3, \dots\}$ of ground instances $P_i = \langle D, I_i \rangle$ over common **first-order STRIPS** domain D
- Each instance $P = \langle D, I \rangle$ consists of:
 - ▷ General (reusable) domain D specified with **action schemas** and **predicates**
 - ▷ Instance information I details **objects**, **init** and **goal** descriptions

Distinction between **general** domain D and **specific** instance $P = \langle D, I \rangle$ important for **reusing** action models, and also for **learning** them

Example (Input): 2-Gripper Problem $P = \langle D, I \rangle$ in PDDL

```
(define (domain gripper)
  (:requirements :typing)
  (:types        room ball gripper)
  (:constants    left right - gripper)
  (:predicates   (at-robot ?r - room) (at ?b - ball ?r - room)
                 (free ?g - gripper) (carry ?o - ball ?g - gripper))

  (:action MOVE
    :parameters  (?from ?to - room)
    :precondition (at-robot ?from)
    :effect       (and (at-robot ?to) (not (at-robot ?from))))

  (:action PICK
    :parameters  (?obj - ball ?room - room ?gripper - gripper)
    :precondition (and (at ?obj ?room) (at-robot ?room) (free ?gripper))
    :effect       (and (carry ?obj ?gripper) (not (at ?obj ?room)) (not (free ?gripper))))

  (:action DROP
    :parameters  (?obj - ball ?room - room ?gripper - gripper)
    :precondition (and (carry ?obj ?gripper) (at-robot ?room))
    :effect       (and (at ?obj ?room) (free ?gripper) (not (carry ?obj ?gripper)))))

(define (problem easy-2balls)
  (:domain gripper)
  (:objects roomA roomB - room B1 B2 - ball)
  (:init (at-robot roomA) (free left) (free right) (at B1 roomA) (at B2 roomA))
  (:goal (and (at B1 roomB) (at B2 roomB))))
```

Relational GNN Architecture for Planning [Ståhlberg et al., 2022-2024]

- Planning state S over STRIPS domain D is a **relational structure**:
 - ▷ Relational symbols given by predicates in D ; **shared** by all such states S
 - ▷ Denotation of predicate p given by ground atoms $p(\bar{o})$ true at S
- Adapt architecture of [Toenshoff et al., 2021] for handling relational structures

Relational GNN (R-GNN) Architecture

Input: Set of ground atoms S (state), and objects O

Output: Embeddings $\mathbf{f}_L(o)$ for each object $o \in O$

1. Initialize $\mathbf{f}_0(o) = 0^k$ for each object $o \in O$
2. **for** $i \in \{0, 1, \dots, L - 1\}$ **do**
3. **for each** atom $q = p(o_1, o_2, \dots, o_m) \in S$ **do**
4. $m_{q,o_j} := [\text{MLP}_p(\mathbf{f}_i(o_1), \mathbf{f}_i(o_2), \dots, \mathbf{f}_i(o_m))]_j$
5. **end for**
6. **for each** object $o \in O$ **do**
7. $\mathbf{f}_{i+1}(o) := \mathbf{f}_i(o) + \text{MLP}_U(\mathbf{f}_i(o), \text{agg}(\{\mathbf{m}_{q,o} \mid o \in q, q \in S\}))$
8. **end for**
9. **end for**

Parameters: embedding dimension k , rounds L , $\{\text{MLP}_p : p \in D\}$, MLP_U , and aggregator

Final Readout, Value Functions, and Greedy Policies



- **Final readout** is **additive readout** that feeds into final MLP:

$$V(S) = \text{MLP}\left(\sum_{o \in O} \mathbf{f}_L(o)\right)$$

- Training minimize **loss** $L(S) = |V^*(S) - V(S)|$ given by **optimal value function** $V^*(\cdot)$ for small tasks in training set
- **Greedy policy** $\pi_V(S)$ chooses action $a = \operatorname{argmin}_{a \in A(S)} 1 + V(S_a)$:
 - ▷ If $V(S) = 0$ for goals, and $V(S) = 1 + \min_a V(S_a)$ for non-goals, π_V is **optimal**
 - ▷ If $V(S) = 0$ for goals, and $V(S) \geq 1 + \min_a V(S_a)$ for non-goals, π_V **solves any state** S where S_a is result of applying action a in state S

Successful approach for GP, but subject to expressivity of GNNs...

Example: Navigation With XY Coordinates




- Navigation in rectangular grid with decoupled coordinates: cells and blocked cells with $\text{CELL}(x, y)$ and $\text{BLOCKED}(x, y)$, position with $\text{AT}(x, y)$, and $\text{ADJ}(i, i + 1)$
- For computing goal distances (ie V^*), cells (x, y) must “communicate” with neighbors (x, y') and (x', y) . In the plain R-GNN, there must be atoms involving $\{x, y, y'\}$ (similarly, $\{x, x', y\}$). No such atoms exists in state S

Expressivity of GNNs

- R-GNNs are instances of (1-)GNNs over undirected graphs
- GNNs compute invariant (resp. equivariant) funcs on graphs (resp. vertices)
- Well-understood **expressivity limitations** in terms of **Weisfeiler-Leman** colorings and **\mathcal{C}_2 logic** (formulas with counting quantifiers, and at most 2 variables)
- Eg, join $W(x, y) = \exists z.[R(x, z) \wedge T(z, y)]$ of relations R and T **cannot be captured!**
- That is, **no GNN can “track” such implicit relation $W(x, y)$** on a graph where red and blue edges stand for R and T respectively
- Can augment expressivity with k -GNNs, $k > 1$, that embed k -tuples of vertices:
 - ▷ Expressivity characterized in terms of k -WL colorings
 - ▷ Either k -OWL (less powerful) or k -FWL (more powerful) versions
 - ▷ Related to, respectively, \mathcal{C}_{k-1} and \mathcal{C}_k logics: counting quant., k variables
 - ▷ **Infeasible by num. objs. in planning problems: $\Theta(N^k)/\Theta(N^{k+1})$ space/time**



Parametric R-GNN $[t]$ Architecture

- Same R-GNN architecture, different encoding of planning states S
- Embedding of all objects pairs, like in 2-GNNs: $\Theta(N^2)$ space
 - ▷ Objects in atoms replaced by pairs: $p(a, b) \rightarrow p(\langle a, a \rangle, \langle a, b \rangle, \langle b, a \rangle, \langle b, b \rangle)$
 - ▷ Predicate arities expanded from k to k^2
- **New composition predicate** $\Delta(\langle x, z \rangle, \langle z, y \rangle, \langle x, y \rangle)$:



- ▷ Set $A_t(S)$ of added Δ -atoms controlled by integer parameter $t \geq 0$
- ▷ $A_0(S) = \{p(\langle w \rangle^2) \mid p(w) \in S\}$ for $\langle w \rangle^2 = \langle (o_1, o_1), \dots, (o_i, o_j), \dots, (o_m, o_m) \rangle$
- ▷ $A_t(S) = A_0(S) \cup \{\Delta(\langle o, o' \rangle, \langle o', o'' \rangle, \langle o, o'' \rangle) \mid \langle o, o' \rangle, \langle o', o'' \rangle \in R_t\}$
- ▷ $\langle o, o' \rangle \in R_t$ iff o and o' in some atom in S ($t=1$), or $\exists o'' [\langle o, o'' \rangle, \langle o'', o' \rangle \in R_{t-1}]$ ($t > 1$)
- $\text{R-GNN}[t](S, O) = \text{R-GNN}(A_t(S), O^2)$
- **Final readout:** $V(S) = \text{MLP}(\sum_{o \in O} \mathbf{f}_L(o, o))$ aggregates $|O|$ embeddings

Example: Navigation With XY Coordinates



- Navigation in rectangular grid with decoupled coordinates: cells and blocked cells with $\text{CELL}(x, y)$ and $\text{BLOCKED}(x, y)$, position with $\text{AT}(x, y)$, and $\text{ADJ}(i, i + 1)$
- After 12 hours of training on 105 random $n \times m$ instances, $mn < 30$, greedy policies achieve coverages of **59.72%**, **80.55%**, and **100%** for R-GNN, R-GNN[0], and R-GNN[1] on instances with **different sets of blocked cells** and $nm \leq 32$
- For computing goal distances (ie V^*), cells (x, y) must “communicate” with neighbors (x, y') and (x', y) . In the plain R-GNN, there must be atoms involving $\{x, y, y'\}$ (similarly, $\{x, x', y\}$). No such atoms exists in S , except in R-GNN[t] where $A_t(S)$ includes $\Delta(\langle x, x' \rangle, \langle x', y \rangle, \langle x, y \rangle)$ and $\Delta(\langle x, y \rangle, \langle y, y' \rangle, \langle x, y' \rangle)$

Experiments: Setup

- A learned value function V for domain D defines a **general policy** π_V that at state S selects an unvisited successor state S' with lowest $V(S')$ value
- We implemented in PyTorch, and trained on Nvidia A10s with 24Gb of memory over 12 hours, using Adam, lr=0.0002, batches of size 16, and no regularization. Embedding dimension of $k = 64$, and $L = 30$ layers were used.
- Standard benchmarks from **International Planning Competition (IPC)**
- For each domain and architecture, 3 models were trained, and best model on validation was selected.
- **Baselines:**
 - ▷ Edge Transformer (ET) [Bergen *et al.*, 2021] designed to do triangulations on graphs
 - ▷ R-GNN₂ that adds all Δ atoms
 - ▷ 2-GNNs that emulates 2-OWL which captures \mathcal{C}_3

Experiments: Results

Domain	Model	Coverage (%)	Plan Length			Domain	Model	Coverage (%)	Plan Length		
			Total	Median	Mean				Total	Median	Mean
Blocks-s	R-GNN	17 / 17 (100 %)	674	38	39	Grid	R-GNN	9 / 20 (45 %)	109	11	12
	R-GNN[0]	17 / 17 (100 %)	670	36	39		R-GNN[0]	12 / 20 (60 %)	177	11	14
	R-GNN[1]	17 / 17 (100 %)	684	36	40		R-GNN[1]	15 / 20 (75 %)	209	13	13
	R-GNN ₂	14 / 17 (82 %)	922	35	65		R-GNN ₂	10 / 20 (50 %)	124	11.5	12
	2-GNN	17 / 17 (100 %)	678	36	39		2-GNN	6 / 20 (30 %)	82	11.5	13
	ET	16 / 17 (94 %)	826	38	51		ET	1 / 20 (5 %)	15	15	15
Blocks-m	R-GNN	22 / 22 (100 %)	868	40	39	Logistics	R-GNN	10 / 20 (50 %)	510	51	51
	R-GNN[0]	22 / 22 (100 %)	830	39	37		R-GNN[0]	9 / 20 (45 %)	439	48	48
	R-GNN[1]	22 / 22 (100 %)	834	39	37		R-GNN[1]	20 / 20 (100 %)	1,057	52	52
	R-GNN ₂	22 / 22 (100 %)	936	39	42		R-GNN ₂	15 / 20 (75 %)	799	52	53
	2-GNN	20 / 22 (91 %)	750	40	37		2-GNN	0 / 20 (0 %)	–	–	–
	ET	18 / 22 (82 %)	966	39	53		ET	0 / 20 (0 %)	–	–	–
Gripper	R-GNN	18 / 18 (100 %)	4,800	231	266	Rovers	R-GNN	9 / 20 (45 %)	2,599	280	288
	R-GNN[0]	18 / 18 (100 %)	1,764	98	98		R-GNN[0]	14 / 20 (70 %)	2,418	153	172
	R-GNN[1]	11 / 18 (61 %)	847	77	77		R-GNN[1]	14 / 20 (70 %)	1,654	55	118
	R-GNN ₂	18 / 18 (100 %)	1,764	98	98		R-GNN ₂	11 / 20 (55 %)	2,225	239	202
	2-GNN	1 / 18 (6 %)	53	53	53		2-GNN	Unsuitable domain: ternary predicates			
	ET	4 / 18 (22 %)	246	61	61		ET				
Miconic	R-GNN	20 / 20 (100 %)	1,342	67	67	Vacuum	R-GNN	20 / 20 (100 %)	4,317	141	215
	R-GNN[0]	20 / 20 (100 %)	1,566	71	78		R-GNN[0]	20 / 20 (100 %)	183	9	9
	R-GNN[1]	20 / 20 (100 %)	2,576	71	128		R-GNN[1]	20 / 20 (100 %)	192	9	9
	R-GNN ₂	20 / 20 (100 %)	1,342	67	67		R-GNN ₂	20 / 20 (100 %)	226	9	11
	2-GNN	12 / 20 (60 %)	649	54.5	54		2-GNN	Unsuitable domain: ternary predicates			
	ET	20 / 20 (100 %)	1,368	68	68		ET				
Visitall	R-GNN	18 / 22 (82 %)	636	29	35	Visitall-xy	R-GNN	5 / 20 (25 %)	893	166	178
	R-GNN[0]	21 / 22 (95 %)	1,128	35	53		R-GNN[0]	15 / 20 (75 %)	1,461	84	97
	R-GNN[1]	22 / 22 (100 %)	886	35	40		R-GNN[1]	20 / 20 (100 %)	1,829	83	91
	R-GNN ₂	20 / 22 (91 %)	739	33	36		R-GNN ₂	19 / 20 (95 %)	2,428	116	127
	2-GNN	18 / 22 (82 %)	626	32	34		2-GNN	12 / 20 (60 %)	1,435	115	119
	ET	18 / 22 (82 %)	670	29	37		ET	3 / 20 (15 %)	455	138	151

Experiments: #Objects in Training / Validation, and Test Sets

Domain	Training / Validation	Test
Blocks	4–9	10–20
Gripper	2–14	16–50
Logistics	2–5 / 3–5	15–19 / 8–11
Miconic	2–20 / 1–10	11–30 / 22–60
Rovers	2–3 / 3–8	3 / 21–39
Vaccum	8–38 / 11–6	40–93 / 6–10
Visitall	1–21	100

Expressivity of the R-GNN $[t]$ Architecture

- The architecture R-GNN $[t]$ has the capability to capture compositions of binary relations that can be expressed in \mathcal{C}_3

Definition (\mathcal{C}_3 -Joins). Let σ be relational language. The class $\mathcal{J}_3 = \mathcal{J}_3[\sigma]$ of relational joins is the **smallest class** of formulas that satisfies:

1. $\{R(x, y), \neg R(x, y)\} \subseteq \mathcal{J}_3$ for relation R in σ ,
2. \mathcal{J}_3 is closed under conjunctions and disjunctions, and
3. $\exists y[\phi(x, y) \wedge \phi(y, z)] \in \mathcal{J}_3$ if $\{\phi(x, y), \phi(y, z)\} \subseteq \mathcal{J}_3$.

Notation $\phi(x, y)$: ϕ is a formula whose free variables are among $\{x, y\}$

Theorem. Let σ be relational language, and let $\mathcal{D} \subseteq \mathcal{J}_3$ be **finite collection** of \mathcal{C}_3 -joins. There is parameter $\langle t, k, L \rangle$, where k is embedding dimension and L is number of layers, and network N in R-GNN $[\sigma, t, k, L]$ that **computes** \mathcal{D}

Conclusions

- Novel parametric architecture R-GNN[t] that **provably** increases the expressivity of the relational R-GNN architecture
- R-GNN[t] embeds all pair of objects but does a bounded number of triangulations, determined by the value of parameter $t \geq 0$
- In benchmarks, a small value of $t = 1$ achieves best results
- Other ways to increase expressivity, like k -GNNs for $k \geq 2$, in either the OWL or FWL setting are infeasible in practice due to high number of objects: $\Theta(N^k)$ space, $\Theta(N^{k+1})$ time
- **Future work:** consider use of indexicals/markers that can be moved around as an alternative to increase the expressivity in GNN architectures for planning

References

- [Bajpai et al., 2018] Bajpai, A. N., Garg, S., et al. (2018). Transfer of deep reactive policies for mdp planning. In *Advances in Neural Information Processing Systems*, pages 10965–10975.
- [Barceló et al., 2020] Barceló, P., Kostylev, E., Monet, M., Pérez, J., Reutter, J., and Silva, J.-P. (2020). The logical expressiveness of graph neural networks. In *Proc. of the 8th International Conference on Learning Representations (ICLR 2020)*. OpenReview.net.
- [Bergen et al., 2021] Bergen, L., O'Donnell, T., and Bahdanau, D. (2021). Systematic generalization with edge transformers. *Advances in Neural Information Processing Systems*, 34:1390–1402.
- [Bonet et al., 2019] Bonet, B., Francès, G., and Geffner, H. (2019). Learning features and abstract actions for computing generalized plans. In *Proc. of the 33rd AAAI Conference on Artificial Intelligence (AAAI 2019)*, pages 2703–2710. AAAI Press.
- [Frances et al., 2021] Frances, G., Bonet, B., and Geffner, H. (2021). Learning general planning policies from small examples without supervision. In *Proc. AAAI*, pages 11801–11808.
- [Grohe, 2021] Grohe, M. (2021). The logic of graph neural networks. In *36th Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*, pages 1–17.
- [Müller et al., 2024] Müller, L., Kusuma, D., Bonet, B., and Morris, C. (2024). Towards principled graph transformers. In *Proc. of the 38th Annual Conference on Neural Information Processing Systems (NeurIPS 2024)*.
- [Ståhlberg et al., 2022a] Ståhlberg, S., Bonet, B., and Geffner, H. (2022a). Learning general optimal policies with graph neural networks: Expressive power, transparency, and limits. In *Proc. ICAPS*, pages 629–637.

- [Ståhlberg et al., 2022b] Ståhlberg, S., Bonet, B., and Geffner, H. (2022b). Learning generalized policies without supervision using GNNs. In *Proc. KR*, pages 474–483.
- [Ståhlberg et al., 2023] Ståhlberg, S., Bonet, B., and Geffner, H. (2023). Learning general policies with policy gradient methods. In *Proceedings of the International Conference on Principles of Knowledge Representation and Reasoning*, pages 647–657.
- [Toenshoff et al., 2021] Toenshoff, J., Ritzert, M., Wolf, H., and Grohe, M. (2021). Graph neural networks for maximum constraint satisfaction. *Frontiers in Artificial Intelligence and Applications*, 3:580607.
- [Toyer et al., 2018] Toyer, S., Trevizan, F., Thiébaux, S., and Xie, L. (2018). Action schema networks: Generalised policies with deep learning. In *AAAI*.