

# Representation Learning for Acting and Planning: A Top Down Approach

Tutorial IJCAI 2022

Blai Bonet

Universitat Pompeu Fabra, Barcelona

Hector Geffner

ICREA & Universitat Pompeu Fabra

Wallenberg Guest Professor, Linköping University

*With inputs from Dominik, Simon, Andrés; RLeap Team (UPF, LiU)*

**Slides** at <https://www.dtic.upf.edu/~hgeffner/tutorial-2022.pdf>



# Bottom-up vs. Top-Down Representation Learning (1)

- Deep learning (DL) and Deep Reinforcement Learning (DRL) have **revolutionised** the landscape of AI, exploiting power of stochastic gradient descent
- Yet DL and DRL struggle with OOD/structural **generalization**
  - ▷ Inductive biases in neural architectures assumed to help but vague, informal
- **Alternative: Language-based representation learning**
  - ▷ Don't choose low-level arch and expect "right representation" to **emerge**
  - ▷ Choose high-level language instead, and learn **representations over language**
- Separation between **what** is to be learned and **how**


## Bottom-up vs. Top-down Representation Learning (2)

- Yoshua Bengio at IJCAI 2021: *System 2 Deep Learning: Higher-Level Cognition, Agency, Out-of-Distribution Generalization and Causality:*

“... **Systematic generalization** hypothesized to arise from efficient factorization of knowledge into **recomposable pieces** corresponding to reusable factors ...”

- Language-based representation learning:
  - ▷ learn the “**recomposable pieces**” in a **language**
  - ▷ recombinations and generalization will follow **semantics**
- Very much in line with **traditional AI**: just *learn from data the representations that have traditionally been crafted by hand*
- **Potential benefits**: meaningful learning bias, semantics, transparency, reasoning

## Example: Minigrid/BabyAI [Chevalier-Boisvert et al., 2019]



- ▶ **Task:** *Pick up grey box behind you, then go to grey key and open door*
- ▶ Red triangle is agent at bottom right. Light-grey is field of view
- ▶ Learn **controller** that accepts **goals** and **obs**, and outputs **action** to do
- ▶ Like a “classical planning problem” **but** state representation **not known**, and goals to be achieved **reactively** (not by planning) with policies that **generalize**

# DRL vs. Language-based Representation Learning

- Surprise is not that DL and DRL methods struggle in Minigrid, but that they manage to generate meaningful behavior at all, given **so little prior knowledge**
- Yet **methodology** largely **ad hoc**: from intuitions to **architectures** and **experiments** using baselines . . .
- From perspective of **language-based representation learning**, **key questions** are:
  - ▷ What are the **domain-independent languages** for representing *dynamics*?
  - ▷ What the **languages** for representing *general reactive policies, subgoals*?
  - ▷ How can **representations** over such languages be **learned**?

# Outline of the Tutorial

- **Background 1:** Classical planning, planning **width**
- **Languages for**
  - ▷ representing general **dynamics**
  - ▷ representing general **policies**
  - ▷ representing general **subgoal structures** (sketches; ‘intrinsic rewards’)
- **Background 2:** Qualitative numerical planning problems (**QNP**s)
- **Learning** representations over these languages:
  - ▷ learning general **dynamics**
  - ▷ learning general **policies**
  - ▷ learning general **subgoal structures**
- **Wrap up; Challenges**

**Copy of these slides** at <https://www.dtic.upf.edu/~hgeffner/tutorial-2022.pdf>

## Outline of the Tutorial (2)

- Tutorial is **not a survey** on learning to act and plan; too much for us; too much for 1:30h
- Focus is on a **coherent** research thread that covers a lot of ground:
  - ▷ **Crisp** and **simple** ideas and formulations for **stating**, **understanding**, and **addressing** key problems
- Many **open problems**; many opportunities for research

# **Background 1:**

## **Classical Planning and Planning Width**



# Background: Model for Classical AI Planning

A (classical) **state model** is a tuple  $\mathcal{S} = \langle S, s_0, S_G, Act, A, f, c \rangle$ :

- finite and discrete **state space**  $S$
- a known **initial state**  $s_0 \in S$
- a set  $S_G \subseteq S$  of **goal states**
- **actions**  $A(s) \subseteq Act$  **applicable** in each  $s \in S$
- a **deterministic state-transition function**  $s' = f(a, s)$  for  $a \in A(s)$
- positive **action costs**  $c(a, s)$ , assumed 1 by default

A **solution** to the model or **plan** is a sequence of applicable actions  $a_0, \dots, a_n$  that maps  $s_0$  into  $S_G$

i.e. there must be state sequence  $s_0, \dots, s_{n+1}$  such that  $a_i \in A(s_i)$ ,  $s_{i+1} = f(a_i, s_i)$ , and  $s_{n+1} \in S_G$

# A Language for Classical Planning: STRIPS

- A (grounded) **problem** in STRIPS is a tuple  $P = \langle F, O, I, G \rangle$ :
  - ▷  $F$  is set of (ground) **atoms**
  - ▷  $O$  is set of (ground) **actions**
  - ▷  $I \subseteq F$  stands for **initial situation**
  - ▷  $G \subseteq F$  stands for **goal situation**
- Actions  $o \in O$  **represented** by
  - ▷ **Add** list  $Add(o) \subseteq F$
  - ▷ **Delete** list  $Del(o) \subseteq F$
  - ▷ **Precondition** list  $Pre(o) \subseteq F$

A **problem**  $P$  in STRIPS defines **state model**  $S(P)$  in compact form . . .

# From Language to Models

STRIPS problem  $P = \langle F, O, I, G \rangle$  determines **state model**  $\mathcal{S}(P)$  where

- the states  $s \in \mathcal{S}$  are collections of atoms from  $F$
- the initial state  $s_0$  is  $I$
- the goal states  $s_G$  are such that  $G \subseteq s_G$
- the actions  $a$  in  $A(s)$  are ops in  $O$  s.t.  $Prec(a) \subseteq s$
- the next state is  $s' = [s \setminus Del(a)] \cup Add(a)$
- action costs  $c(a, s)$  are all 1

Common approach for solving  $P$  is using **path-finding/heuristic search** algorithms over **graph** defined by  $\mathcal{S}(P)$  where nodes are the states  $s$ , and edges  $(s, s')$  are state transitions caused by an action  $a$ ; i.e.,  $s' = f(a, s)$  and  $a \in A(s)$

The **source** node is the initial state  $s_0$ , and the **targets** are the goal states  $s_G$

## Background: Width and Width-based Algorithms

- IW(1) is a **breadth-first search** that **prunes** states  $s$  that don't make a **feature** true for first time in the search, given **set of Boolean features**  $F$ 
  - ▷ In **classical planning**,  $F$  is the set of (ground) atoms in problem
- IW( $k$ ) is IW(1) but over set  $F^k$  made up of conjunctions of  $k$  features from  $F$
- **Alternatively**, IW( $k$ ) is a breadth-first search that prunes  $s$  if **novelty**( $s$ )  $> k$
- **IW** runs IW(1), IW(2), . . . , IW( $k$ ) sequentially until problem solved or  $k = N$
- IW is blind like DFS and BFS but diff **enumeration**; uses **state structure**
- IW( $k$ ) expands up to  $N^k$  nodes and runs in **polytime**  $\exp(2k - 1)$

# Planning for \*Atomic Goals\* with IW(1) and IW(2)

#	Domain	I	IW(1)	IW(2)	Neither
1.	8puzzle	400	55%	45%	0%
2.	Barman	232	9%	0%	91%
3.	Blocks World	598	26%	74%	0%
4.	Cybersecure	86	65%	0%	35%
...	...	...	...	...	...
22.	Pegsol	964	92%	8%	0%
23.	Pipes-NonTan	259	44%	56%	0%
24.	Pipes-Tan	369	59%	37%	3%
25.	PSRsmall	316	92%	0%	8%
26.	Rovers	488	47%	53%	0%
27.	Satellite	308	11%	89%	0%
28.	Scanalyzer	624	100%	0%	0%
...	...	...	...	...	...
33.	Transport	330	0%	100%	0%
34.	Trucks	345	0%	100%	0%
35.	Visitall	21,859	100%	0%	0%
36.	Woodworking	1659	100%	0%	0%
37.	Zeno	219	21%	79%	0%
<b>Total/Avg</b>		<b>37,921</b>	<b>37.0%</b>	<b>51.3%</b>	<b>11.7%</b>

**88.3%** of the 37,921 instances solved by IW(1) or IW(2) [Lipovetzky and G., 2012]

# Performance of IW is No Accident: Theory

- **Width** of  $P$ ,  $w(P)$ , is min  $k$  for which there is a sequence of **subgoals** (atom tuples)  $t_0, t_1, \dots, t_n$ ,  $|t_i| \leq k$  such that:
  - ▷  $t_0$  is true in the initial situation
  - ▷ the optimal plans for  $t_n$  are optimal plans for  $P$
  - ▷ all **optimal plans for  $t_i$  can be extended into optimal plans for  $t_{i+1}$  by adding a single action**
- Also  $w(P) = 0$  if goal reachable in 0 or 1 step;  $w(P) = N + 1$  if no solution, where  $N$  is number of atoms in  $P$ .
- **Theorem:** If  $w(P) = k$ , then IW( $k$ ) solves  $P$  optimally in  $\exp(2k - 1)$  time
- **Theorem:** Domains like Blocks, Logistics, Gripper, ... have all **width 2** independent of problem **size** provided that goals are **single atoms**

## Practical Variations of IW

**SIW:** Serialized iterated width [Lipovetzky and G., 2012]

- Use IW greedily to decrease **number of unachieved goals**  $\#g$ ; assumes conjunctive top goal (simple goal serialization)

**BFWS:** Best-first guided by **novelty measure**  $w_{\langle \#g, \#c \rangle}$  and  $\#g$

- $\text{BFWS}(f_5)$ : back-end of state-of-the-art Dual-BFWS,  $\#c$  from relaxed plans
- $k\text{-BFWS}(f_5)$ : **poltytime** variant of  $\text{BFWS}(f_5)$  used as front-end of Dual-BFWS
- $\text{BFWS}(R)$ : version that does not use **action structure**; just **PDDL simulator**

[Lipovetzky and G., 2017; Francès *et al.*, 2017]

# Understanding Width: Test Your Knowledge!

How to **prove** in standard encodings that:

- Blocks world instances with goal  $clear(x)$  or  $hold(x)$  have **width 1**
- Delivery instances with goal  $hold(x)$  or  $AgentAt(y)$  have **width 1**
- Blocks world instances with goal  $on(x, y)$  have **width 2**
- Delivery instances with goal  $PkgAt(x, y)$  have **width 2**
- Blocks and Delivery with **arbitrary conjunctive goals** have **no bounded width**

**Delivery** is simplified LOGISTICS: agent in grid, picking up and dropping pkgs

For **proving**  $w(G) \leq k$ :

- **Necessary 1:** If  $a_1, \dots, a_n$  is optimal plan for goal  $G$ , each **prefix**  $a_1, \dots, a_i$  must be optimal plan for some  $t_i$ ,  $|t_i| \leq k$
- **Necessary 2:** For these  $t_i$ 's, **all** optimal plans for  $t_i$  **extend** into optimal plans for  $t_{i+1}$ .



## Part II: Languages

- Language for expressing **dynamics**
- Language for expressing **general policies**
- Language for expressing **general subgoal structures**

# Language for Expressing Dynamics: First-Order STRIPS

Problems specified as **instances**  $P = \langle D, I \rangle$  of **general** planning domain:

- **Domain**  $D$  specified in terms of **action schemas** and **predicates**
- **Instance** is  $P = \langle D, I \rangle$  where  $I$  details **objects, init, goal**

Distinction between **general** domain  $D$  and **specific** instance  $P = \langle D, I \rangle$  important for **reusing** action models, and also for **learning** them:

- Learning  $P_i = \langle D, I_i \rangle$  implies learning  $D$  that **generalizes** to other instances

In RL and DRL, there is no notion of **domain**: generalization to other “instances” analyzed **experimentally**; closest things are “procedurally generated instances,” and “probability distribution over tasks”

## Example: 2-Gripper Problem $P = \langle D, I \rangle$ in PDDL

```
(define (domain gripper)
  (:requirements :typing)
  (:types room ball gripper)
  (:constants left right - gripper)
  (:predicates (at-robot ?r - room)(at ?b - ball ?r - room)(free ?g - gripper)
    (carry ?o - ball ?g - gripper))
  (:action move
    :parameters (?from ?to - room)
    :precondition (at-robot ?from)
    :effect (and (at-robot ?to) (not (at-robot ?from))))
  (:action pick
    :parameters (?obj - ball ?room - room ?gripper - gripper)
    :precondition (and (at ?obj ?room) (at-robot ?room) (free ?gripper))
    :effect (and (carry ?obj ?gripper) (not (at ?obj ?room)) (not (free ?gripper))))
  (:action drop
    :parameters (?obj - ball ?room - room ?gripper - gripper)
    :precondition (and (carry ?obj ?gripper) (at-robot ?room))
    :effect (and (at ?obj ?room) (free ?gripper) (not (carry ?obj ?gripper)))))

(define (problem gripper2)
  (:domain gripper)
  (:objects roomA roomB - room Ball1 Ball2 - ball)
  (:init (at-robot roomA) (free left) (free right) (at Ball1 roomA)(at Ball2 roomA))
  (:goal (and (at Ball1 roomB) (at Ball2 roomB))))
```

# Preview: Learning Dynamics in Lifted STRIPS

- Planning problem  $P_i = \langle D, I_i \rangle$  defines unique **state graph**  $G(P_i)$
- Learning as **inverse problem**: from graphs  $G_1, \dots, G_k$ , learn  $D, I_i$ :

Given graphs  $G_1, \dots, G_k$ , find **simplest** instances  $P_i = \langle D, I_i \rangle$  such that graphs  $G_i$  and  $G(P_i)$  are isomorphic,  $i = 1, \dots, k$ .

- Problem cast and solved as combinatorial optimization task [B. and G., 2020]
- **Complexity** of  $D$  determined by  $\#$  and arities of action schemas and predicates
- **Variations**: missing edges, noisy observations [Rodriguez *et al.*, 2021a]
- **Related**
  - ▷ Learning schemas from **ground traces** [Cresswell *et al.*, 2013]
  - ▷ Deep learning of action schemas from images via **autoencoders** [Asai, 2019]
  - ▷ Learning prop. action models from **options** [Konidaris *et al.*, 2018]
  - ▷ Most work on **learning action models** assumes **domain predicates** known

## Second Task: General Policies

- **General policy** represents strategy for solving **multiple** domain instances **reactively**; i.e., without having to search or plan
  - ▷ E.g., policy for achieving  $on(x, y)$ ; **any** # of blocks, **any** configuration
- What are good **languages** for expressing such policies?
- Number of works have addressed the problem [Khardon 1999; Martin and G., 2004; Fern *et al.*, 2006; Srivastava *et al.*, 2011; Hu and De Giacomo, 2011]
- **Subtlety**: set of (ground) actions change from instance to instance with objects

**Learning general policies** also a key goal in (Deep) RL

# General Policies: A Language [B. and G., 2018]

- **General policies** are given by **rules**  $C \mapsto E$  over set  $\Phi$  of **features**
- **Features**  $f$  are state functions that have a well-defined value  $f(s)$  on every reachable state of any instance of the domain
  - ▷ **Boolean** features  $p$ :  $p(s)$  is true or false
  - ▷ **Numerical** features  $n$ :  $n(s)$  is non-negative integer

Computation of feature values assumed to be “cheap”: features assumed to have **linear** number of values at most, computable in **linear** time (in  $|P|$ ).

## Example: General Policy for $clear(x)$

- **Features**  $\Phi = \{H, n\}$ : 'holding' and 'number of blocks above  $x$ '
- **Policy**  $\pi$  for class  $\mathcal{Q}$  of Block problems with goal  $clear(x)$  given by two rules:

$$\{\neg H, n > 0\} \mapsto \{H, n\downarrow\} \quad ; \quad \{H, n > 0\} \mapsto \{\neg H\}$$

### Meaning:

- if  $\neg H$  &  $n > 0$ , move to successor state where  $H$  holds and  $n$  **decreases**
- if  $H$  &  $n > 0$ , move to successor state where  $\neg H$  holds,  $n$  **doesn't change**

# Language and Semantics of General Policies: Definitions

- **Policy rules**  $C \mapsto E$  over set  $\Phi$  of Boolean and numerical **features**  $p, n$ :
  - ▷ *Boolean conditions* in  $C$ :  $p, \neg p, n = 0, n > 0$
  - ▷ *qualitative effects* in  $E$ :  $p, \neg p, p?, n\downarrow, n\uparrow, n?$
- **State transition**  $(s, s')$  **satisfies** rule  $C \mapsto E$  if
  - ▷  $f(s)$  makes body  $C$  true
  - ▷ change from  $f(s)$  to  $f(s')$  satisfies  $E$
- A **policy**  $\pi$  for class  $\mathcal{Q}$  of problems  $P$  is given by policy rules  $C \mapsto E$ 
  - ▷ *Transition*  $(s, s')$  in  $P$  compatible with  $\pi$  if  $(s, s')$  satisfies a policy rule
  - ▷ *Trajectory*  $s_0, s_1, \dots$  compatible if  $s_0$  of  $P$  and transitions compatible with  $\pi$
- $\pi$  **solves**  $P$  if all max trajectories compatible with  $\pi$  reach goal of  $P$
- $\pi$  **solves** collection of problems  $\mathcal{Q}$  if it solves each  $P \in \mathcal{Q}$



## Example: Delivery

- Pick packages spread in  $n \times m$  grid, one by one, to target location
- **Features**  $\Phi = \{H, p, t, n\}$ : hold, dist. to nearest pkg & target, # undelivered
- Policy  $\pi$  that solves class  $\mathcal{Q}_D$ : **any** # of pkgs and distribution, **any** grid size

$\{\neg H, p > 0\} \mapsto \{p\downarrow, t?\}$	go to nearest package
$\{\neg H, p = 0\} \mapsto \{H, p?\}$	pick it up
$\{H, t > 0\} \mapsto \{t\downarrow, p?\}$	go to target cell
$\{H, t = 0\} \mapsto \{\neg H, n\downarrow, p?\}$	drop package

# General Policies: Three Questions

1. How to **prove** that general policy solves potentially infinite class of instances  $Q$ ?
2. How to **learn** policies (and the features involved) to solve  $Q$ ?
3. How to **learn** policies that are **guaranteed** to solve infinite  $Q$ ?

We consider idea of **learning** first and move then to 1. Not much to say about 3.

## Preview: Learning General Policies

Given a known domain  $D$ , training instances  $P_1, \dots, P_k$ , over  $D$ , and a **finite pool of domain features**  $\mathcal{F}$ , each with a cost, find the cheapest policy  $\pi$  over  $\mathcal{F}$  such that  $\pi$  solves all  $P_i$ ,  $i = 1, \dots, k$

- Problem cast and solved as **combinatorial opt. task** [Francès *et al.*, 2021]
- Pool of **features**  $\mathcal{F}$  generated from domain predicates using **2-variable** (description) logic grammar; feature cost given by syntax tree size
- **Deep learning** approaches [Toyer *et al.*, 2018; Garg *et al.*, 2020] do not need  $\mathcal{F}$  but not 100% correct in general
- Recent DL approach also avoids  $\mathcal{F}$  and nearly 100% correct when **2-variable logic** features suffice; exploits relation between **GNNs** and 2-variable logic [Ståhlberg *et al.*, 2022a and 2022b]

# Proving that a General Policy Solves Class of Instances $\mathcal{Q}$

How to **prove** that this policy  $\pi$  achieves  $clear(x)$  in all Block problems?

$$\{\neg H, n > 0\} \mapsto \{H, n \downarrow\} \quad ; \quad \{H, n > 0\} \mapsto \{\neg H\}$$

- **Soundness:** policy  $\pi$  applies in every **non-goal** state  $s$ 
  - ▷ for any such  $s$ , there is  $(s, s')$  compatible with  $\pi$
- **Acyclicity:** no sequence of transitions  $(s_i, s_{i+1})$  compatible with  $\pi$  **cycle**

**Theorem:** If  $\pi$  is sound and acyclic in  $\mathcal{Q}$ , and no dead-ends,  $\pi$  solves  $\mathcal{Q}$

**Exercise:** Show that policy for  $clear(x)$  is **sound** and **acyclic** in *Blocks*

# Acyclicity, Termination, and QNPs

- **Termination:** criterion that ensures that policy is **acyclic** over **any** domain
- A policy  $\pi$  is **terminating** if for all infinite trajectories  $s_0, \dots, s_i, \dots$  compatible with  $\pi$ , there is a **numerical feature**  $n$  such that:
  - ▷  $n$  is **decremented** in some recurrent transition  $(s, s')$ ; i.e.,  $n(s') < n(s)$
  - ▷  $n$  is **not incremented** in any recurrent transition  $(s, s')$ ; i.e.,  $n(s') \not> n(s)$
- Every such trajectory deemed **impossible** or **unfair** ( $n$  can't decrement below 0), thus if  $\pi$  terminates,  $\pi$ -trajectories **terminate**
- **Termination** notion is from **QNPs**; verifiable in time  $O(2^{|\Phi|})$  by SIEVE algorithm [Srivastava *et al.*, 2011], where  $\Phi$  is set of features involved in the policy

More about QNPs later on . . .

## Third Task: Subgoal Structure

**Subgoal structure** important in planning and RL (“intrinsic rewards”, hierarchies)

**Sketches** powerful language for expressing subgoal structure [B. and G., 2021]

- **Goal serialization** and **full policies** expressible as sketches
- **Semantics** in terms of **subgoals to be achieved**; not so with HTNs, LTL
- Sketches **split** problems into **subproblems**

If **subproblems** have a **bounded width**, problems solved in **polytime**

## Example: Sketches for Delivery

- **Width=0** Sketch (full policy)

$\{\neg H, p > 0\} \mapsto \{p\downarrow, t?\}$	go to nearest package
$\{\neg H, p = 0\} \mapsto \{H, p?\}$	pick it up
$\{H, t > 0\} \mapsto \{t\downarrow, p?\}$	go to target cell
$\{H, t = 0\} \mapsto \{\neg H, n\downarrow, p?\}$	drop package

- **Width=2** Sketch:

$\{n > 0\} \mapsto \{n\downarrow\}$	deliver package
-------------------------------------	-----------------

- **Width=1** Sketch:

$\{\neg H\} \mapsto \{H\}$	go and pick package
$\{H\} \mapsto \{\neg H, n\downarrow\}$	go and deliver package

**Features:** Holding ( $H$ ); Dist. to nearest Pkg ( $p$ ), Target ( $t$ ); # Undeliv Pkgs ( $n$ )

# Syntax and Semantics of Sketch Rules

- **Syntax:** For Boolean and numerical **features**  $p$  and  $n$ :
  - ▷  $p, \neg p, n > 0, n = 0$  can appear in  $C$
  - ▷  $p, \neg p, n\uparrow, n\downarrow, n?$  can appear in  $E$
- **Semantics:** **State pair**  $(s, s')$  **satisfies** sketch rule  $C \mapsto E$  if
  - ▷  $f(s)$  satisfies  $C$
  - ▷  $(f(s), f(s'))$  satisfies  $E$

Syntax of sketches and policies the **same**, and so with semantics, **except** that  $(s, s')$  is not a **1-step state transition** necessarily

**Interpretation:** When in state  $s$ , the **set of subgoal states**  $G_R(s)$  to aim at is:

$$G_R(s) = \{ s' \mid (s, s') \text{ satisfies sketch rule or } s' \text{ is goal} \}$$



# Sketch Width

- Sketch  $R$  **splits** problems  $P$  in  $\mathcal{Q}$  into collection of **subproblems**  $P[s, G_R(s)]$ :
  - ▷ **Initial state**  $s$ : reachable state  $s$  in  $P$
  - ▷ **(Sub) goal states**  $G_R(s) = \{ s' \mid (s, s') \text{ satisfies sketch rule or } s' \text{ is goal} \}$
- **Width of sketch**  $R$  over  $\mathcal{Q} = \max_{s, P \in \mathcal{Q}} \text{width}(P[s, G_R(s)])$ 
  - ▷ *for definition in presence of **dead-ends**, see refs*

**Theorem:** Any  $P$  in  $\mathcal{Q}$  is **solvable** in  $O(b \cdot N^{|\Phi|+2k-1})$  time by  $\text{SIW}_R$  algorithm if sketch  $R$  is **terminating** and has **width** over  $\mathcal{Q}$  bounded by  $k$  [B. and G., 2021]

▷  $N$ : Number of atoms in problem  $P$  ;  $\Phi$ : Set of features in sketch

$\text{SIW}_R$  is like SIW but **subgoal** to achieve next given by sketch

▷ SIW is  $\text{SIW}_R$  with sketch  $R$  with **single rule**:  $\{\#g > 0\} \mapsto \{\#g \downarrow\}$

## Another Example: IPC Grid [Drexler et al., 2021]

This sketch is **terminating** and has **width** 1 for IPC domain Grid (pick and deliver keys spread in grid where cells can be locked and opened with other keys):

- **Sketch:**

- ▷  $r_1 : \{l > 0\} \mapsto \{l\downarrow, k?, o?, t?\}$  (if locked cells, unlock them)
- ▷  $r_2 : \{l = 0, k > 0\} \mapsto \{k\downarrow, o?, t?\}$  (else, place keys in targets)
- ▷  $r_3 : \{l > 0, \neg o\} \mapsto \{o, t?\}$  (if locked cells, pick key to open locked cell)
- ▷  $r_4 : \{l = 0, \neg t\} \mapsto \{o?, t\}$  (if all locks open and misplaced keys, pick up such key)

- **Features:**

- ▷  $l$  is the number of unlocked grid cells
- ▷  $k$  is the number of misplaced keys
- ▷  $o$  is true iff robot holds key for which there is a closed lock
- ▷  $t$  is true iff robot holds key that must be placed at some target grid cell

## Preview: Learning Sketches [Drexler et al., 2022]

Given a known domain  $D$ , training instances  $P_1, \dots, P_n$ , and non-negative integer  $k$ , find simplest sketch  $R$  over a pool of features  $\mathcal{F}$  such that

- Subproblems induced by  $R$  on each  $P_i$  have all **width bounded** by  $k$ ,
- Sketch  $R$  is **terminating**

Possibly first approach for **learning subgoal structure** based on crisp **principles**

**Many threads that come together:**

- Planning **width**
- Language of **general policies**
- Termination notion from **QNP**s
- Semantics of **sketches**

## Exercise: Test Your Knowledge! (Not trivial)

In the 1985 AIJ paper, *Macro-Operators: A Weak Method for Learning*, Rich Korf provides **macro-tables** for puzzles like Rubik Cube, 8-puzzle, and other hard puzzles that encode **policies**  $\pi(s)$  for solving them from any initial state

- Can these compact policies be replaced by even more compact **sketches** of **bounded width**?
- Can these sketches be **general**? That is, applicable to Rubik cubes and  $n$ -sliding puzzles of **different sizes**?
- Can such sketches be **learned** with current method? Expressivity? Scalability? Other methods?

## **Background 2:**

# **Qualitative Numerical Planning Problems (QNPs)**

# Language for QNPs

- Language for planning involving **propositional** and **numerical variables**
- QNPs [Srivastava *et al.* 2011] different than **numerical planning**:
  - ▷ Numerical vars in QNPs are non-negative, **real-valued**
  - ▷ **Effects** on numerical variables: just **qualitative** increments/decrements
  - ▷ **Numerical literals**: whether variable is **zero** or **positive** only
- These differences make plan-existence for QNPs **decidable**
- QNPs provide language for **general policies and sketches**:
  - ▷ QNP actions similar to policy/sketch rules but **features** replaced by **variables**
- We follow [B. and G., 2020b]

## Syntax for QNPs

A **qualitative numerical problem (QNP)** is tuple  $Q = \langle F, V, I, O, G \rangle$ :

- $F$  and  $V$  are sets of propositional and numerical **variables** (not features!)
- $I$  and  $G$  denote initial and goal states
- $O$ : actions  $a$  with precs, and prop. and numeric effects  $Pre(a)$ ,  $Eff(a)$ ,  $N(a)$ :
  - ▷  $F$ -literals may appear in  $I$ ,  $G$ ,  $Pre(a)$  and  $Eff(a)$
  - ▷  $V$ -literals may appear in  $I$ ,  $G$  and  $Pre(a)$
  - ▷  $N(a)$  can only have expressions of the form  $X\uparrow$  and  $X\downarrow$  for var  $X$  in  $V$
- $V$ -literal is either  $X = 0$  or  $X > 0$  for variable  $X$  in  $V$
- **Example:** QNP  $Q_{clear} = \langle \{H\}, \{n\}, I, O, G \rangle$ 
  - ▷  $I = \{n > 0, \neg H\}$
  - ▷  $G = \{n = 0\}$
  - ▷  $O = \{a, b\}$  where  $a = \{\neg H, n > 0\} \mapsto \{H, n\downarrow\}$  and  $b = \{H\} \mapsto \{\neg H\}$
- QNP actions like policy rules above but  $H$  and  $n$  not features but **variables**

# Semantics and Solutions of QNPs

- Policy  $\pi$  for a QNP is partial map from state  $s$  into actions such that:
  - ▷  $\pi(s) = \pi(s')$  if  $s$  and  $s'$  **qualitatively similar**: same  $F$  and  $V$  true literals
- $\pi$  **solves** QNP if all maximal **QNP-fair**  $\pi$ -trajectories reach the goal
  - ▷ **QNP fairness**: trajectory **unfair** if numerical variable **decremented** infinite number of times and **incremented** finite number of times.

**Theorem** [Srivastava *et al.*, 2011]:  $\pi$  solves QNP  $Q$  iff  $\pi$  is **strong cyclic solution** of the **FOND** problem  $T_D(Q)$  obtained from  $Q$  that **terminates**

- $T_D(Q)$  replaces **numerical**  $X$  by **Boolean** variable “ $X > 0$ ” (“ $X = 0$ ” is negative literal)
- **Qualitative effects**  $X \uparrow$  replaced by **effect**  $X > 0$
- **Qualitative effects**  $X \downarrow$  replaced by **non-deterministic effect** “ $X > 0 \mid X = 0$ ”
- **Strong-cyclic**: every reachable state is connected to goal state by  $\pi$

Polytime reduction from QNPs to FOND, but more complex than  $T_D$  [B. and G., 2020b]



# Termination, Sieve Algorithm [Srivastava et al., 2011]

Policy for QNP  $Q$  **terminates** if no infinite **QNP-fair**  $\pi$ -trajectories

SIEVE provides **sound** and **complete** polynomial termination test

- State  $s$  **terminates** if either
  - ▷ there is no cycle on state  $s$ , or
  - ▷ every cycle on  $s$  contains a state  $s'$  that terminates, or
  - ▷  $\pi(s)$  decrements a variable  $X$ , and every cycle on  $s$  that contains a state  $s'$  such that  $\pi(s')$  increments  $X$ , contains another state  $s''$  that terminates
- Policy  $\pi$  terminates iff every state reached by  $\pi$  terminates

Recent FOND<sup>+</sup> planner handles strong FOND, strong cyclic FOND, QNPs, and hybrids by stating **fairness assumptions** explicitly [Rodriguez *et al.* 2021b]

## Part III: Learning Dynamics, Policies, Sketches

- Learning **action models**:

Given graphs  $G_1, \dots, G_k$ , find **simplest** instances  $P_i = \langle D, I_i \rangle$  such that graphs  $G_i$  and  $G(P_i)$  are **isomorphic**,  $i = 1, \dots, k$ .

- Learning **general policies**:

Given known domain  $D$ , training instances  $P_1, \dots, P_k$ , over  $D$ , and **finite pool of domain features**  $\mathcal{F}$ , each with a cost, find the cheapest policy  $\pi$  over  $\mathcal{F}$  such that  $\pi$  **solves all**  $P_i$ ,  $i = 1, \dots, k$

- Learning **sketches**:

Given known domain  $D$ , training instances  $P_1, \dots, P_n$ , and non-negative integer  $k$ , find simplest sketch  $R$  over a pool of features  $\mathcal{F}$  such that

- ▷ Subproblems induced by  $R$  on each  $P_i$  have all **width bounded** by  $k$ ,
- ▷ Sketch  $R$  is **terminating**

# Learning Action Models: Encoding [Rodriguez et al., 2021a]

- Construct **answer set program**, bounding number of objects, preds, and action/pred. arities:
  - ▷ **Given**  $G_1, \dots, G_n$  as input graphs over **black-box states**, with edge labels,
  - ▷ **Check** whether there is STRIPS model  $D$  and instances  $I_1, \dots, I_n$  such that graphs  $G(P_i)$  and  $G_i$  are **isomorphic**,  $i = 1, \dots, n$ , where  $P_i = \langle D, I_i \rangle$
  - ▷ **Optimize**: sum of action and predicate arities, etc
- **(Basic) choice variables**:
  - ▷ Lifted atom is pair  $(P, T)$  where  $P$  is int and  $T$  is tuple of ints
  - ▷  $\text{prec}(A, (P, T), V)$  and  $\text{eff}(A, (P, T), V)$  (lifted atoms in precs/effects)
  - ▷  $\text{p\_arity}(P, N)$  and  $\text{a\_arity}(A, N)$  (arities for predicate and action)
  - ▷  $\text{val}(S, (P, O), V)$  where  $O$  is tuple of objs and  $V$  is 0/1 (value of ground atoms at states)
  - ▷  $\text{appl}(A, O, S)$  and  $\text{next}(A, O, S, T)$  (ground action  $A(O)$  appl/assigned to  $(S, T)$ )
- **(Basic) constraints**:
  - ▷  $\text{:- state}(S), \text{state}(T), S < T, \text{val}(T, (P, O), V) : \text{val}(S, (P, O), V).$  (diff. states)
  - ▷  $\{ \text{next}(A, O, S, T) : \text{label}((S, T), A) \} = 1 \text{ :- } \text{appl}(A, O, S).$  (assign edges to actions)
  - ▷  $\text{:- state}(S), \text{action}(A), N = \{ \text{label}((S, T), A) \}, \{ \text{appl}(A, O, S) \} \neq N.$  (matching)
  - ▷ Compliance of precs/effects of assigned grounded actions to edges
- CLINGO program  $\sim 400$  lines [Rodriguez et al. 2021a]; more complex in SAT [B. and G., 2020a]

# Learning General Policies: Encoding [Francès et al., 2021]

- **Input** is set of transitions  $\mathcal{S}$  from small instances, pool of features  $\mathcal{F}$ , **parameter** (int)  $\delta$
- **Output** is policy: rules obtained from **selected features** and (“good”) **transitions**
- **Combinatorial opt. task**  $T(\mathcal{S}, \mathcal{F}, \delta)$ : Solve constraints minimizing **feature complexity**
- **Choice variables:**
  - ▷ `select(F)` (features that define rules)
  - ▷ `good(S,T)` (transition (S,T) is “compatible” with policy)
  - ▷ `V(S,N)` (distance from S to goal is N)
- **Constraints:**
  - ▷ `1 { good(S,T) } :- state(S), not terminal(S).` (good transitions at non-terminals)
  - ▷ `:- good(S,T), deadend(T).` (no good tr. reaches dead-end T)
  - ▷ `1 { select(F) : diff(F,S,T) } :- goal(S), not goal(T).` (distinguish goals)
  - ▷ `{ V(S,D) :  $V^*(S) \leq D \leq \delta V^*(S)$  } = 1 :- state(S).` (set distances)
  - ▷ `:- good(S,T), V(S,D1), V(T,D2), D2 <= D1.` (distances avoid cycles)
  - ▷ `1 { select(F) : diff(F,S1,T1,S2,T2) } :- good(S1,T1), not good(S2,T2).` (distinguish good/bad transitions)

where `diff/3` and `diff/5` computed from pool at **pre-processing**

# Learning General Sketches: Encoding [Drexler et al., 2022]

- **Input:** transitions  $\mathcal{S}$  in small instances, pool  $\mathcal{F}$ , width bound  $k$ , max # sketch rules  $m$
- **Output:** sketch of **width**  $\leq k$ , **acyclic** in given instances, with up to  $m$  rules
- **Combinatorial opt. task**  $T(\mathcal{S}, \mathcal{F}, k, m)$ : solve constraints min **complexity** of selected features
- **(Basic) variables:**
  - ▷ `rule(I)` (sketch rule I)
  - ▷ `select(F)` (features that define sketch rules)
  - ▷ `cond(I,F,V)` and `eff(I,F,E)` (conditions and effects for rule I)
  - ▷ `subgoal(S,T)` (tuple T of width  $k$  is subgoal for S)
  - ▷ **(Implied)** `subgoal(S1,T,S2)` (subgoal T for S1 may lead to S2)
  - ▷ **(Implied)** `satis(S1,S2,I)` (pair (S1,S2) satisfies rule I)
- **(Basic) constraints:**
  - ▷ **Well formed rules:** atoms `cond/3` and `eff/3` are consistent and imply `select(F)`
  - ▷ `1 { subgoal(S,T) : tuple(T) } :- state(S), not goal(S).` (width  $k$  subgoal for S)
  - ▷ `subgoal(S1,T,S2) :- subgoal(S1,T), found(S1,T,S2).` (subgoal T may lead to S2)
  - ▷ `:- subgoal(S1,T,S2), not satis(S1,S2,I) : rule(I).` ((S1,S2) satisfies some rule)
  - ▷ `:- satis(S1,S2,I), not subgoal(S1,T) : d(S1,T) < d(S1,S2).` (dead-end S2 is farther)
  - ▷ `:- satis(S1,S2,I), not subgoal(S1,T) : d(S1,T) ≤ d(S1,S2).` (subgoals optimal)
  - ▷ Collection of rules is **terminating** (approx'ed by testing acyclicity)

# About the Pool of Features $\mathcal{F}$ [B. et al., 2019]

- **Description logic grammar** allows generation of **concepts** and **roles** from **domain predicates**
- Complexity of concept/role given by **size of its syntax tree**
- Pool  $\mathcal{F}$  obtained from concepts of complexity bounded by parameter
- Denotation of concept  $C$  in state  $s$  is **subset of objects**
- Each concept  $C$  defines num and Bool features  $n_C(s) = |C(s)|$ ;  $p_C(s) = \top$  iff  $|C(s)| > 0$
- Grammar:
  - ▷ Primitive:  $C_p$  given by unary predicates  $p$  and unary “goal predicates”  $p_G$
  - ▷ Universal:  $C_u$  contains all objects
  - ▷ Nominals:  $C_a = \{a\}$  for constants/parameter  $a$
  - ▷ Negation:  $\neg C$  contains  $C_u \setminus C$
  - ▷ Intersection:  $C \sqcap C'$
  - ▷ Quantified:  $\exists R.C = \{x : \exists y[R(x, y) \wedge C(y)]\}$  and  $\forall R.C = \{x : \forall y[R(x, y) \wedge C(y)]\}$
  - ▷ Roles (for binary predicate  $p$ ):  $R_p$ ,  $R_p^{-1}$ ,  $R_p^+$ , and  $[R_p^{-1}]^+$
- Additional **distance features**  $dist(C_1, R, C_2)$  for concepts  $C_1$  and  $C_2$  and role  $R$  that evaluates to  $d$  in state  $s$  iff minimum  $R$ -distance between object in  $C_1$  to object in  $C_2$  is  $d$

# General Policies By Deep Learning [Ståhlberg et al., 2022a,b]

- Exploits correspondence between **graph neural networks (GNNs)** and **two-variable logic  $\mathcal{C}_2$**  to learn policy **without requiring** pool of  $\mathcal{C}_2$  features  $\mathcal{F}$
- **Value function  $V$**  learned that yields general policy  $\pi_V$  **greedy** in  $V$
- For **generalization**, based on GNN arch. for MaxCSP( $\Gamma$ ) [Toenshoff et al., 2021]
  - ▷ **Input** given by the states  $s$  extended with “goal predicates”  $p_G$
  - ▷ **Output**  $V(s)$  is non-linear aggregation of object embeddings
  - ▷ **Loss:**  $|V^*(s) - V(s)|$  for supervised learning of optimal policies
  - ▷ **Loss:**  $\max\{0, [1 + \min_{s' \in N(s)} V(s')] - V(s)\}$  unsupervised/non-optimal
- Nearly **as good as** policies based on **explicit pool  $\mathcal{F}$  of  $\mathcal{C}_2$  features**
- Complexity of “latent features” not explicitly bounded

# GNN Architecture [Ståhlberg et al., 2022a,b]

---

**Algorithm 1:** GNN maps state  $s$  into scalar  $V(s)$

---

**Input:** State  $s$ : set of atoms true in  $s$ , set of objects

**Output:**  $V(s)$

```
1  $f_0(o) \sim \mathbf{0}^{k/2} \mathcal{N}(0, 1)^{k/2}$  for each object  $o \in s$ ;  
2 for  $i \in \{0, \dots, L - 1\}$  do  
3   for each atom  $q := p(o_1, \dots, o_m)$  true in  $s$  do  
4      $m_{q,o} := [\mathbf{MLP}_p(f_i(o_1), \dots, f_i(o_m))]_j$ ;  
5   for each  $o$  in  $s$  do  
6      $f_{i+1}(o) := \mathbf{MLP}_U(f_i(o), \text{agg}(\{m_{q,o} | o \in q\}))$ ;  
    $//$  Aggregate, update embeddings  
7  $V := \mathbf{MLP}_2(\sum_{o \in s} \mathbf{MLP}_1(f_L(o)))$ 
```

---



# Wrap Up: Representation Learning for Acting and Planning

- **Background 1:** Classical planning, planning **width**
- **Languages** for
  - ▷ representing general dynamics
  - ▷ representing general policies
  - ▷ representing general subgoal structures (sketches; ‘intrinsic rewards’)
- **Background 2:** Qualitative numerical planning problems (**QNP**s)
- **Learning** representations over these languages:
  - ▷ learning general dynamics
  - ▷ learning general policies
  - ▷ learning general subgoal structures
- **Wrap up; Challenges**

# Wrap Up

- To learn representations that generalize due to structure, don't play with low-level neural architecture; choose suitable (domain-independent) **target language** and learn representations over it:
  - ▷ generalization
  - ▷ transparency
  - ▷ powerful, meaningful bias
  - ▷ distinction between **what** and **how**
- Examples of learning language-based representations to **act** and **plan**:
  - ▷ general action **dynamics**
  - ▷ general **policies**
  - ▷ general **subgoal structures** (sketches)

# Challenges: Language-based Representation Learning

- Scalability of combinatorial optimization approaches
- Use of deep learning (learning lifted dynamics, policies, sketches).
- Alternative target languages for learning (e.g., vs. lifted STRIPS)
- Continuous domains, space, time
- Stochastic and non-deterministic domains
- States in the input: black-box, parsed images, images, videos
- Grounded vs. ungrounded representations
- Learning and reusing “skills”, hierarchies
- . . .

<https://www.dtic.upf.edu/~hgeffner/tutorial-2022.pdf>

Plenty to do; if seriously interested, reach us

# References

- [Asai, 2019] Asai, M. (2019). Unsupervised grounding of plannable first-order logic representation from images. In *Proc. ICAPS*.
- [Bonet et al., 2019] Bonet, B., Francès, G., and Geffner, H. (2019). Learning features and abstract actions for computing generalized plans. In *Proc. AAAI*, pages 2703–2710.
- [Bonet and Geffner, 2018] Bonet, B. and Geffner, H. (2018). Features, projections, and representation change for generalized planning. In *Proc. IJCAI*, pages 4667–4673.
- [Bonet and Geffner, 2020a] Bonet, B. and Geffner, H. (2020a). Learning first-order symbolic representations for planning from the structure of the state space. In *Proc. ECAI*.
- [Bonet and Geffner, 2020b] Bonet, B. and Geffner, H. (2020b). Qualitative numeric planning: Reductions and complexity. *Journal of AI Research*, 69:923–961.
- [Bonet and Geffner, 2021] Bonet, B. and Geffner, H. (2021). General policies, representations, and planning width. In *Proc. AAAI*, pages 11764–11773.
- [Chevalier-Boisvert et al., 2019] Chevalier-Boisvert, M., Bahdanau, D., Lahlou, S., Willems, L., Saharia, C., Nguyen, T. H., and Bengio, Y. (2019). Babyai: A platform to study the sample efficiency of grounded language learning. In *ICLR*.
- [Cresswell et al., 2013] Cresswell, S. N., McCluskey, T. L., and West, M. M. (2013). Acquiring planning domain models using LOCM. *The Knowledge Engineering Review*, 28(2):195–213.
- [Drexler et al., 2021] Drexler, D., Seipp, J., and Geffner, H. (2021). Expressing and exploiting the common subgoal structure of classical planning domains using sketches. In *Proc. KR*, pages 258–268.
- [Drexler et al., 2022] Drexler, D., Seipp, J., and Geffner, H. (2022). Learning sketches for decomposing planning problems into subproblems of bounded width. In *Proc. ICAPS*.
- [Fern et al., 2006] Fern, A., Yoon, S., and Givan, R. (2006). Approximate policy iteration with a policy language bias: Solving relational markov decision processes. *Journal of Artificial Intelligence Research*, 25:75–118.

- [Francès et al., 2021] Francès, G., Bonet, B., and Geffner, H. (2021). Learning general planning policies from small examples without supervision. In *Proc. AAAI*, pages 11801–11808.
- [Garg et al., 2020] Garg, S., Bajpai, A., and Mausam (2020). Symbolic network: generalized neural policies for relational mdps. In *International Conference on Machine Learning*, pages 3397–3407.
- [Geffner and Bonet, 2013] Geffner, H. and Bonet, B. (2013). *A Concise Introduction to Models and Methods for Automated Planning*. Morgan & Claypool Publishers.
- [Hu and De Giacomo, 2011] Hu, Y. and De Giacomo, G. (2011). Generalized planning: Synthesizing plans that work for multiple environments. In *Proc. IJCAI*, pages 918–923.
- [Kharden, 1999] Kharden, R. (1999). Learning action strategies for planning domains. *Artificial Intelligence*, 113:125–148.
- [Konidaris et al., 2018] Konidaris, G., Kaelbling, L. P., and Lozano-Perez, T. (2018). From skills to symbols: Learning symbolic representations for abstract high-level planning. *Journal of Artificial Intelligence Research*, 61:215–289.
- [Lipovetzky and Geffner, 2012] Lipovetzky, N. and Geffner, H. (2012). Width and serialization of classical planning problems. In *Proc. ECAI*, pages 540–545.
- [Lipovetzky and Geffner, 2017a] Lipovetzky, N. and Geffner, H. (2017a). Best-first width search: Exploration and exploitation in classical planning. In *Proc. AAAI*.
- [Lipovetzky and Geffner, 2017b] Lipovetzky, N. and Geffner, H. (2017b). A polynomial planning algorithm that beats lama and ff. *Proc. ICAPS*.
- [Martín and Geffner, 2004] Martín, M. and Geffner, H. (2004). Learning generalized policies from planning examples using concept languages. *Applied Intelligence*, 20(1):9–19.
- [Rodriguez et al., 2021a] Rodriguez, I. D., Bonet, B., Romero, J., and Geffner, H. (2021a). Learning first-order representations for planning from black-box states: New results. In *KR*. arXiv preprint arXiv:2105.10830.
- [Rodriguez et al., 2021b] Rodriguez, I. D., Bonet, B., Sardina, S., , and Geffner, H. (2021b). Flexible fond planning with explicit fairness assumptions. In *Proc. ICAPS*, pages 290–298.
- [Srivastava et al., 2011] Srivastava, S., Zilberstein, S., Immerman, N., and Geffner, H. (2011). Qualitative numeric planning. In *AAAI*.

- [Ståhlberg et al., 2022a] Ståhlberg, S., Bonet, B., and Geffner, H. (2022a). Learning general optimal policies with graph neural networks: Expressive power, transparency, and limits. In *Proc. ICAPS*.
- [Ståhlberg et al., 2022b] Ståhlberg, S., Bonet, B., and Geffner, H. (2022b). Learning generalized policies without supervision using gnns. In *Proc. KR*.
- [Toenshoff et al., 2021] Toenshoff, J., Ritzert, M., Wolf, H., and Grohe, M. (2021). Graph neural networks for maximum constraint satisfaction. *Frontiers in artificial intelligence*, 3:98.
- [Toyer et al., 2018] Toyer, S., Trevizan, F., Thiébaux, S., and Xie, L. (2018). Action schema networks: Generalised policies with deep learning. In *AAAI*.