

# Learning Depth-First Search: A Unified Approach to Heuristic Search in Deterministic and Non-Deterministic Settings

Blai Bonet

Dept. de Computación  
Universidad Simón Bolívar  
Caracas, Venezuela

Héctor Geffner

Dept. de Tecnología  
ICREA & Universitat Pompeu Fabra  
Barcelona, Spain

# Motivation

- **Heuristic search** methods can be **efficient** but lack common foundation: IDA\*, AO\*, Alpha-Beta, ...
- **Dynamic programming** methods such as Value Iteration are **general** but not as efficient:
  - Single algorithm for wide range of models: Det, MDPs, Games, AND/OR, ...
  - yet VI is exhaustive
- This work aims to bring these two types of methods together to obtain:
  - efficiency, generality, and understanding!

# Result

- A simple algorithm, **Learning Depth-First Search** (LDFS), capable of solving a **wide range** of deterministic and non-deterministic models; based on three ideas
  - Depth-First Search
  - Lower bounds
  - Learning
- For some models, LDFS **reduces to state-of-the-art algorithms**:
  - **Deterministic Models**: LDFS = IDA\* w/ transposition tables
  - **Game Trees**: (Bounded) LDFS = Alpha-Beta w/ null windows (MTD) [Plaat et. al, 1996]
  - On others, like AND/OR and MDPs, LDFS yields **new algorithms**

# Basic Intuitions: IDA\*

---

## IDA\*

- Performs iterative Depth-First Searches with certain *bound*
- Prune action  $a$  in node  $n$  leading to node  $n'$  when

$$g(n) + c(a, n) + h(n') > bound$$

# Basic Intuitions: IDA\*

## IDA\*


- Performs iterative Depth-First Searches with certain *bound*
- Prune action  $a$  in node  $n$  leading to node  $n'$  when

$$g(n) + c(a, n) + h(n') > bound$$


## IDA\* w/ Transposition Table (Cost Revisions [Reinefeld & Marsland, 1994])

- As IDA\*: performs iteratives DFS's with certain *bound*
- Upon backtracking from node  $n$ , **revise** heuristic value  $h(n)$  (stored in TT) to the new lower bound


# Illustration: IDA\* w/ Transposition Table



# Illustration: IDA\* w/ Transposition Table




# Illustration: IDA\* w/ Transposition Table







# Illustration: IDA\* w/ Transposition Table




# Illustration: IDA\* w/ Transposition Table




# Illustration: IDA\* w/ Transposition Table




# Illustration: IDA\* w/ Transposition Table




# Illustration: IDA\* w/ Transposition Table




# Illustration: IDA\* w/ Transposition Table



# Illustration: IDA\* w/ Transposition Table



# Illustration: IDA\* w/ Transposition Table





# Monotone Heuristics and IDA\*

- Monotonicity is a property of  $h$  that says

$$f(n) \leq f(n')$$

for each node  $n$  and successor  $n'$ ; i.e. **non-decreasing  $f$ -values along paths**

- With monotone  $h$ , IDA\* has **important properties**:
  - The bound for each iteration equals  $h(n_0)$
  - An iteration either finds solution or increases  $h(n_0)$  which is next bound
  - The revision of the heuristic (TT) is just

$$h(n) := \min_{a \in A(n)} c(n, a) + h(n')$$

- Path  $(n_0, a_0, n_1, \dots, a_i, n_{i+1})$  is **transversed** by IDA\* (w/ TT and Monotone  $h$ ) iff

$$bound = f(n_0) \leq f(n_1) \leq f(n_2) \leq \dots \leq f(n_i) \leq bound$$

$$f(n_0) = f(n_1) = f(n_2) = \dots = f(n_i)$$

# IDA\* + TT and Monotone $h$ : Reformulated (Generalized)

- If  $Q_h(a, n) = c(a, n) + h(n')$ , the algorithm can be expressed as iterations that:
  - Starting from  $n_0$ , perform DFS along actions  $a$  such that

$$h(n) = Q_h(a, n)$$

- Backtrack at tip nodes  $n$  (i.e. with no such  $a$ 's), restoring consistency of  $h(n)$ :

$$h(n) := \min_{a \in A(n)} Q_h(a, n)$$

# IDA\* + TT and Monotone $h$ : Reformulated (Generalized)

- If  $Q_h(a, n) = c(a, n) + h(n')$ , the algorithm can be expressed as iterations that:
  - Starting from  $n_0$ , perform DFS along actions  $a$  such that

$$h(n) = Q_h(a, n)$$

- Backtrack at tip nodes  $n$  (i.e. with no such  $a$ 's), restoring consistency of  $h(n)$ :

$$h(n) := \min_{a \in A(n)} Q_h(a, n)$$

- **Good News:** reformulation is very **general**; other models can be solved **efficiently** by suitable choice of  $Q_h(a, n)$ :

$$Q_h(a, n) = c(a, n) + \max_{n'} h(n') \quad \text{for MAX AND/OR GRAPHS}$$

$$Q_h(a, n) = c(a, n) + \sum_{n'} P_a(n'|n) h(n') \quad \text{for MDPs}$$

$$Q_h(a, n) = \max_{n'} h(n') \quad \text{for GAME TREES}$$

...

...

# IDA\* + TT and Monotone $h$ : Reformulated (Generalized)

- If  $Q_h(a, n) = c(a, n) + h(n')$ , the algorithm can be expressed as iterations that:
  - Starting from  $n_0$ , perform DFS along actions  $a$  such that

$$h(n) = Q_h(a, n)$$

- Backtrack at tip nodes  $n$  (i.e. with no such  $a$ 's), restoring consistency of  $h(n)$ :

$$h(n) := \min_{a \in A(n)} Q_h(a, n)$$

- **Good News:** reformulation is very **general**; other models can be solved **efficiently** by suitable choice of  $Q_h(a, n)$ :

$$Q_h(a, n) = c(a, n) + \max_{n'} h(n') \quad \text{for MAX AND/OR GRAPHS}$$

$$Q_h(a, n) = c(a, n) + \sum_{n'} P_a(n'|n) h(n') \quad \text{for MDPs}$$

$$Q_h(a, n) = \max_{n'} h(n') \quad \text{for GAME TREES}$$

...

...

- We call this algorithm **LDFS** for **Learning Depth-First Search**

# LDFS: The Code

```
LDFS-DRIVER( $n_0$ )
begin
  repeat  $solved := \text{LDFS}(n_0)$  until  $solved$ 
  return  $(V, \pi)$ 
end

LDFS( $n$ )
begin
  if  $n$  is SOLVED or terminal then
    if  $n$  is terminal then  $V(n) := c_T(n)$ 
    Mark  $n$  as solved return true

  flag := false % EXPANSION
  foreach  $a \in A(n)$  do
    if  $Q_V(a, n) > V(n)$  then continue
    flag := true
    foreach  $n' \in F(a, n)$  do
      flag := LDFS( $s'$ ) & flag % Recursion
      if  $\neg flag$  then break
    if flag then break

  if flag then
     $\pi(n) := a$ 
    Mark  $n$  as SOLVED % LABELING
  else
     $V(n) := \min_{a \in A(n)} Q_V(a, n)$  % UPDATE
  return flag
end
```

## Rest of the Talk: Outline

---

- What are Q-factors and where they come from?
- Why LDFS works?
- Properties and relation to other algorithms
- Extensions and empirical results for MDPs

## A bit of Background: Models

- a discrete and finite states space  $S$ ,
- an initial state  $s_0 \in S$ ,
- a non-empty set of terminal states  $S_T \subseteq S$ ,
- actions  $A(s) \subseteq A$  applicable in each non-terminal state,
- a function that maps states and actions into sets of states  $F(a, s) \subseteq S$ ,
- action costs  $c(a, s)$  for non-terminal states  $s$ , and
- terminal costs  $c_T(s)$  for terminal states.

DETERMINISTIC:  $|F(a, s)| = 1$  (OR Graphs),

NON-DETERMINISTIC:  $|F(a, s)| \geq 1$  (AND/OR graphs),

MDPs: probabilities  $P_a(s'|s)$  for  $s' \in F(s, a)$  that add up to 1

...

# Solutions

- (Optimal) Solutions can be expressed in terms of value function  $V$  satisfying **Bellman** equation:

$$V(s) = \begin{cases} c_T(s) & \text{if } s \text{ is terminal} \\ \min_{a \in A(s)} Q_V(a, s) & \text{otherwise} \end{cases}$$

where  $Q_V(a, s)$  stands for the cost-to-go value defined as:

$Q_V(a, s) = c(a, s) + V(s'), s' \in F(a, s)$	for OR GRAPHS
$Q_V(a, s) = c(a, s) + \max_{s' \in F(a, s)} V(s')$	for MAX AND/OR GRAPHS
$Q_V(a, s) = c(a, s) + \sum_{s' \in F(a, s)} V(s')$	for ADD AND/OR GRAPHS
$Q_V(a, s) = c(a, s) + \sum_{s' \in F(a, s)} P_a(s' s) V(s')$	for MDPs
$Q_V(a, s) = \max_{s' \in F(a, s)} V(s')$	for GAME TREES



# Solutions

- (Optimal) Solutions can be expressed in terms of value function  $V$  satisfying **Bellman** equation:

$$V(s) = \begin{cases} c_T(s) & \text{if } s \text{ is terminal} \\ \min_{a \in A(s)} Q_V(a, s) & \text{otherwise} \end{cases}$$

where  $Q_V(a, s)$  stands for the cost-to-go value defined as:

$Q_V(a, s) = c(a, s) + V(s'), s' \in F(a, s)$	for OR GRAPHS
$Q_V(a, s) = c(a, s) + \max_{s' \in F(a, s)} V(s')$	for MAX AND/OR GRAPHS
$Q_V(a, s) = c(a, s) + \sum_{s' \in F(a, s)} V(s')$	for ADD AND/OR GRAPHS
$Q_V(a, s) = c(a, s) + \sum_{s' \in F(a, s)} P_a(s' s) V(s')$	for MDPs
$Q_V(a, s) = \max_{s' \in F(a, s)} V(s')$	for GAME TREES

- An **optimal policy** can be recovered from the solution of Bellman equation as:

$$\pi(s) = \operatorname{argmin}_{a \in A(s)} Q_V(a, s)$$

# LDFS: Learning Depth-First Search

- Assuming monotone and admissible value function  $V$  (i.e.  $h$ ):
  - Start from  $n_0$  and perform DFS along actions such that

$$V(n) = Q_V(a, n)$$

- Backtrack when there is no such action and update  $V(n)$  to

$$V(n) := \min_{a \in A(n)} Q_V(a, n)$$

- LDFS solves all models above, **except MDPs with cyclic solutions**
- LDFS is **equivalent** to IDA\* w/ TT on Deterministic models (OR graphs)

# Value Iteration Algorithm

- There is an algorithm that is almost as general, and even simpler: **Value Iteration**
- Value Iteration **doesn't search, just makes updates**:
  - Iterate until convergence:

$$\text{For all node } n \text{ do: } V(n) := \min_{a \in A(n)} Q_V(a, n)$$

- VI is pretty good when **all states fit in memory** (e.g. around  $10^6$  states)

## Why is VI less general and less effective than LDFS?

- It doesn't work in the presence of dead-ends (states with  $V(s) = \infty$ )

## Why is VI less general and less effective than LDFS?

- It doesn't work in the presence of dead-ends (states with  $V(s) = \infty$ )
- It doesn't exploit LBs/Heuristic information; **it is an exhaustive method**

# Why is VI less general and less effective than LDFS?

- It doesn't work in the presence of dead-ends (states with  $V(s) = \infty$ )
- It doesn't exploit LBs/Heuristic information; **it is an exhaustive method**
- For example, if  $Q_V(a, n) > V^*(n)$ , LDFS will never consider action  $a$  at  $n$ ;
- E.g. IDA\* never explores child  $n'$  of  $n$  if  $Q_h(a, n) > h(n)$  and  $h$  is monotone

# Find-and-Revise: An Abstraction that Searches and Updates

- Find-and-Revise is a theoretical model for analysis; defined in terms of
  - Greedy Graph  $G_V$ : contains nodes  $n$  reachable from  $n_0$  by applying **greedy** actions  $a$ ; i.e. those with  $Q_V(a, n) = \min_a Q_V(a, n)$
  - Consistent nodes: those such that  $V(n) = \min_a Q_V(a, n)$

# Find-and-Revise: An Abstraction that Searches and Updates

- Find-and-Revise is a theoretical model for analysis; defined in terms of
  - Greedy Graph  $G_V$ : contains nodes  $n$  reachable from  $n_0$  by applying **greedy** actions  $a$ ; i.e. those with  $Q_V(a, n) = \min_a Q_V(a, n)$
  - Consistent nodes: those such that  $V(n) = \min_a Q_V(a, n)$
- Find-and-Revise iterates:
  - Find an inconsistent node  $n$  in greedy graph  $G_V$
  - Update (revise)  $V$  at node  $n$until no such state exists



# Find-and-Revise: An Abstraction that Searches and Updates

- Find-and-Revise is a theoretical model for analysis; defined in terms of
  - Greedy Graph  $G_V$ : contains nodes  $n$  reachable from  $n_0$  by applying **greedy** actions  $a$ ; i.e. those with  $Q_V(a, n) = \min_a Q_V(a, n)$
  - Consistent nodes: those such that  $V(n) = \min_a Q_V(a, n)$
- Find-and-Revise iterates:
  - Find an inconsistent node  $n$  in greedy graph  $G_V$
  - Update (revise)  $V$  at node  $n$until no such state exists
- **THM:** Find-and-Revise solves all models if initial  $V$  is admissible and monotone

# LDFS is an instance of Find-and-Revise!

LDFS is a Find-and-Revise that:

- Finds with a DFS search that backtracks upon inconsistent states
- Upon backtracking updates inconsistent states and ancestors
- Keeps track of SOLVED states to avoid re-exploration (**labeling**)

# Some Properties of LDFS

## Additive Models (e.g. OR graphs, Additive AND/OR, ...)

- Each iteration of LDFS either increases the value of  $s_0$  or labels  $s_0$  as solved
- Hence, number of iterations bounded by  $V^*(s_0) - h(s_0)$
- Does not deal with probabilistic models with cyclic solutions (MDPs)

# Some Properties of LDFS

## Additive Models (e.g. OR graphs, Additive AND/OR, ...)

- Each iteration of LDFS either increases the value of  $s_0$  or labels  $s_0$  as solved
- Hence, number of iterations bounded by  $V^*(s_0) - h(s_0)$
- Does not deal with probabilistic models with cyclic solutions (MDPs)

## Max Models (e.g. Max AND/OR, Game Trees, ...)

- An iteration of LDFS may no increase the value of  $s_0$  neither label it
- Yet a simple variation, called **Bounded LDFS**, restores such property
- **Bounded LDFS = Alpha-Beta w/ null windows** (MTD) [Plaat et. al, 1996]

## Variations Needed for Cyclic MDPs

- Cannot do the bottom-up labeling because of cycles

## Variations Needed for Cyclic MDPs

- Cannot do the bottom-up labeling because of cycles
- In MDPs, the value function converges to  $V^*$  **asymptotically**, thus need to replace consistency

$$\min_a Q_V(a, n) - V(n) = 0$$

with  $\varepsilon$ -consistency:

$$\min_a Q_V(a, n) - V(n) \leq \varepsilon$$

# Variations Needed for Cyclic MDPs

- Cannot do the bottom-up labeling because of cycles
- In MDPs, the value function converges to  $V^*$  **asymptotically**, thus need to replace consistency

$$\min_a Q_V(a, n) - V(n) = 0$$

with  $\varepsilon$ -consistency:

$$\min_a Q_V(a, n) - V(n) \leq \varepsilon$$

- IDA\* might perform an exponential number of iterations in problems with real costs, and so do LDFS in MDPs

# Variations Needed for Cyclic MDPs

- Cannot do the bottom-up labeling because of cycles
- In MDPs, the value function converges to  $V^*$  **asymptotically**, thus need to replace consistency

$$\min_a Q_V(a, n) - V(n) = 0$$

with  $\varepsilon$ -consistency:

$$\min_a Q_V(a, n) - V(n) \leq \varepsilon$$

- IDA\* might perform an exponential number of iterations in problems with real costs, and so do LDFS in MDPs
- The resulting algorithm is called LDFS+



# A Bit of Empirical Results

- Four domains: noisy 8-puzzle, racetracks, rooms, tree
- Algorithms: VI, LRTDP, ILAO, HDP, LDFS+
- Two (monotone) heuristics: zero, min-min relaxation

algorithm	small	big	bigger	ring-1	ring-2	ring-3	ring-4	ring-5	ring-6
$ S $	9,394	22,532	51,941	429	1,301	5,949	33,243	94,396	352,150
$V^*(s_0)$	14.459	26.134	50.570	7.498	10.636	13.093	18.530	24.949	31.142
$h_{min-min}(s_0)$	11	18	37	6	9	11	15	20	25
$VI(h_{min-min})$	1.080	3.824	14.761	0.022	0.105	0.611	5.198	23.168	197.964
$LRTDP(h_{min-min})$	0.369	3.169	12.492	0.006	0.027	0.138	2.173	15.361	243.130
$ILAO(h_{min-min})$	0.813	4.739	20.190	0.008	0.034	0.463	11.428	37.598	—
$HDP(h_{min-min})$	0.468	5.357	30.174	0.007	0.034	0.180	2.159	11.473	153.150
$LDFS+(h_{min-min})$	<b>0.196</b>	<b>1.077</b>	<b>4.542</b>	<b>0.003</b>	<b>0.014</b>	<b>0.083</b>	<b>1.022</b>	<b>4.892</b>	<b>80.068</b>
$VI(h=0)$	1.501	5.289	21.701	0.027	0.124	0.774	7.281	34.501	354.917
$LRTDP(h=0)$	0.880	6.232	29.836	<b>0.012</b>	0.109	0.356	6.005	171.829	—
$ILAO(h=0)$	2.430	14.200	54.208	0.024	0.109	0.908	11.863	71.103	—
$HDP(h=0)$	2.440	30.955	174.698	0.032	0.149	0.927	11.957	96.398	—
$LDFS+(h=0)$	<b>0.792</b>	<b>3.417</b>	<b>16.080</b>	0.013	<b>0.057</b>	<b>0.353</b>	<b>4.390</b>	<b>24.732</b>	<b>310.019</b>

# Summary

- A simple algorithm, **Learning Depth-First Search** (LDFS), capable of solving a **wide range** of deterministic and non-deterministic models; based on three ideas
  - Depth-First Search
  - Lower bounds
  - Learning
- For some models, LDFS **reduces to state-of-the-art algorithms**:
  - **Deterministic Models**: LDFS = IDA\* w/ transposition tables
  - **Game Trees**: (Bounded) LDFS = Alpha-Beta w/ null windows (MTD) [Plaat et. al, 1996]
  - On others, like AND/OR and MDPs, LDFS yields **new algorithms**
- Competitive results for LDFS+ on **Markov Decision Processes**