

Representation Learning for Acting and Planning: A Top Down Approach

Tutorial ICAPS 2022

Blai Bonet

Universitat Pompeu Fabra, Barcelona

Hector Geffner

ICREA & Universitat Pompeu Fabra

Wallenberg Guest Professor, Linköping University

With inputs from Dominik, Simon, Andrés; RLeap Team (UPF, LiU)



Bottom-up vs. Top-Down Representation Learning (1)

- Deep learning (DL) and Deep Reinforcement Learning (DRL) have **revolutionised** the landscape of AI, exploiting power of stochastic gradient descent
- Yet DL and DRL struggle with OOD/structural **generalization**
 - ▷ Inductive biases in neural architectures assumed to help but vague, informal
- **Alternative: Language-based representation learning**
 - ▷ Don't choose low-level arch and expect "right representation" to **emerge**
 - ▷ Choose high-level language instead, and learn **representations over language**
- Separation between **what** is to be learned and **how**


Bottom-up vs. Top-down Representation Learning (2)

- Yoshua Bengio at IJCAI 2021: *System 2 Deep Learning: Higher-Level Cognition, Agency, Out-of-Distribution Generalization and Causality:*

“... **Systematic generalization** hypothesized to arise from efficient factorization of knowledge into **recomposable pieces** corresponding to reusable factors ...”

- Language-based representation learning:
 - ▷ learn the “**recomposable pieces**” in a **language**
 - ▷ recombinations and generalization will follow **semantics**
- Very much in line with **traditional AI**: just *learn from data the representations that have traditionally been crafted by hand*
- **Potential benefits**: meaningful learning bias, semantics, transparency, reasoning

Example: Minigrid/BabyAI [Chevalier-Boisvert et al., 2019]



- ▶ **Task:** *Pick up grey box behind you, then go to grey key and open door*
- ▶ Red triangle is agent at bottom right. Light-grey is field of view
- ▶ Learn **controller** that accepts **goals** and **obs**, and outputs **action** to do
- ▶ Like a “classical planning problem” **but** state representation **not known**, and goals to be achieved **reactively** (not by planning) with policies that **generalize**

DRL vs. Language-based Representation Learning

- Surprise is not that DL and DRL methods struggle in Minigrid, but that they manage to generate meaningful behavior at all, given **so little prior knowledge**
- Yet **methodology** largely **ad hoc**: from intuitions to **architectures** and **experiments** using baselines . . .
- From perspective of **language-based representation learning**, **key questions** are:
 - ▷ What are the **domain-independent languages** for representing *dynamics*?
 - ▷ What the **languages** for representing *general reactive policies, subgoals*?
 - ▷ How can **representations** over such languages be **learned**?

Outline of the Tutorial

- **Background 1:** Classical planning, planning **width**
- **Languages for**
 - ▷ representing general **dynamics**
 - ▷ representing general **policies**
 - ▷ representing general **subgoal structures** (sketches; ‘intrinsic rewards’)
- **Background 2:** Qualitative numerical planning problems (**QNP**s)
- **Learning** representations over these languages:
 - ▷ learning general **dynamics**
 - ▷ learning general **policies**
 - ▷ learning general **subgoal structures**
- **Wrap up; Challenges**

Copy of these slides at <https://www.dtic.upf.edu/~hgeffner/tutorial-2022.pdf>

Outline of the Tutorial (2)

- Tutorial is **not a survey** on learning to act and plan; too much for us; too much for 1:30h
- Focus is on a **coherent** research thread that covers a lot of ground:
 - ▷ **Crisp** and **simple** ideas and formulations for **stating**, **understanding**, and **addressing** key problems
- Many **open problems**; many opportunities for research

Background 1:

Classical Planning and Planning Width

Background: Model for Classical AI Planning

A (classical) **state model** is a tuple $\mathcal{S} = \langle S, s_0, S_G, Act, A, f, c \rangle$:

- finite and discrete **state space** S
- a known **initial state** $s_0 \in S$
- a set $S_G \subseteq S$ of **goal states**
- **actions** $A(s) \subseteq Act$ **applicable** in each $s \in S$
- a **deterministic state-transition function** $s' = f(a, s)$ for $a \in A(s)$
- positive **action costs** $c(a, s)$, assumed 1 by default

A **solution** to the model or **plan** is a sequence of applicable actions a_0, \dots, a_n that maps s_0 into S_G

i.e. there must be state sequence s_0, \dots, s_{n+1} such that $a_i \in A(s_i)$, $s_{i+1} = f(a_i, s_i)$, and $s_{n+1} \in S_G$

A Language for Classical Planning: STRIPS

- A (grounded) **problem** in STRIPS is a tuple $P = \langle F, O, I, G \rangle$:
 - ▷ F is set of (ground) **atoms**
 - ▷ O is set of (ground) **actions**
 - ▷ $I \subseteq F$ stands for **initial situation**
 - ▷ $G \subseteq F$ stands for **goal situation**
- Actions $o \in O$ **represented** by
 - ▷ **Add** list $Add(o) \subseteq F$
 - ▷ **Delete** list $Del(o) \subseteq F$
 - ▷ **Precondition** list $Pre(o) \subseteq F$

A **problem** P in STRIPS defines **state model** $S(P)$ in compact form . . .

From Language to Models

STRIPS problem $P = \langle F, O, I, G \rangle$ determines **state model** $\mathcal{S}(P)$ where

- the states $s \in \mathcal{S}$ are collections of atoms from F
- the initial state s_0 is I
- the goal states s_G are such that $G \subseteq s_G$
- the actions a in $A(s)$ are ops in O s.t. $Prec(a) \subseteq s$
- the next state is $s' = [s \setminus Del(a)] \cup Add(a)$
- action costs $c(a, s)$ are all 1

Common approach for solving P is using **path-finding/heuristic search** algorithms over **graph** defined by $\mathcal{S}(P)$ where nodes are the states s , and edges (s, s') are state transitions caused by an action a ; i.e., $s' = f(a, s)$ and $a \in A(s)$

The **source** node is the initial state s_0 , and the **targets** are the goal states s_G

Background: Width and Width-based Algorithms

- IW(1) is a **breadth-first search** that **prunes** states s that don't make a **feature** true for first time in the search, given **set of Boolean features** F
 - ▷ In **classical planning**, F is the set of (ground) atoms in problem
- IW(k) is IW(1) but over set F^k made up of conjunctions of k features from F
- **Alternatively**, IW(k) is a breadth-first search that prunes s if **novelty**(s) $> k$
- **IW** runs IW(1), IW(2), . . . , IW(k) sequentially until problem solved or $k = N$
- IW is blind like DFS and BFS but diff **enumeration**; uses **state structure**
- IW(k) expands up to N^k nodes and runs in **polytime** $\exp(2k - 1)$

Planning for *Atomic Goals* with IW(1) and IW(2)

#	Domain	I	IW(1)	IW(2)	Neither
1.	8puzzle	400	55%	45%	0%
2.	Barman	232	9%	0%	91%
3.	Blocks World	598	26%	74%	0%
4.	Cybersecure	86	65%	0%	35%
...
22.	Pegsol	964	92%	8%	0%
23.	Pipes-NonTan	259	44%	56%	0%
24.	Pipes-Tan	369	59%	37%	3%
25.	PSRsmall	316	92%	0%	8%
26.	Rovers	488	47%	53%	0%
27.	Satellite	308	11%	89%	0%
28.	Scanalyzer	624	100%	0%	0%
...
33.	Transport	330	0%	100%	0%
34.	Trucks	345	0%	100%	0%
35.	Visitall	21,859	100%	0%	0%
36.	Woodworking	1659	100%	0%	0%
37.	Zeno	219	21%	79%	0%
Total/Avg		37,921	37.0%	51.3%	11.7%

88.3% of the 37,921 instances solved by IW(1) or IW(2) [Lipovetzky and G., 2012]

Performance of IW is No Accident: Theory

- **Width** of P , $w(P)$, is min k for which there is a sequence of **subgoals** (atom tuples) t_0, t_1, \dots, t_n , $|t_i| \leq k$ such that:
 - ▷ t_0 is true in the initial situation
 - ▷ the optimal plans for t_n are optimal plans for P
 - ▷ all **optimal plans for t_i can be extended into optimal plans for t_{i+1} by adding a single action**
- Also $w(P) = 0$ if goal reachable in 0 or 1 step; $w(P) = N + 1$ if no solution, where N is number of atoms in P .
- **Theorem:** If $w(P) = k$, then IW(k) solves P optimally in $\exp(2k - 1)$ time
- **Theorem:** Domains like Blocks, Logistics, Gripper, ... have all **width 2** independent of problem **size** provided that goals are **single atoms**

Practical Variations of IW

SIW: Serialized iterated width [Lipovetzky and G., 2012]

- Use IW greedily to decrease **number of unachieved goals** $\#g$; assumes conjunctive top goal (simple goal serialization)

BFWS: Best-first guided by **novelty measure** $w_{\langle \#g, \#c \rangle}$ and $\#g$

- $\text{BFWS}(f_5)$: back-end of state-of-the-art Dual-BFWS, $\#c$ from relaxed plans
- $k\text{-BFWS}(f_5)$: **poltytime** variant of $\text{BFWS}(f_5)$ used as front-end of Dual-BFWS
- $\text{BFWS}(R)$: version that does not use **action structure**; just **PDDL simulator**

[Lipovetzky and G., 2017; Francès *et al.*, 2017]

Understanding Width: Test Your Knowledge!

How to **prove** in standard encodings that:

- Blocks world instances with goal $clear(x)$ or $hold(x)$ have **width 1**
- Delivery instances with goal $hold(x)$ or $AgentAt(y)$ have **width 1**
- Blocks world instances with goal $on(x, y)$ have **width 2**
- Delivery instances with goal $PkgAt(x, y)$ have **width 2**
- Blocks and Delivery with **arbitrary conjunctive goals** have **no bounded width**

Delivery is simplified LOGISTICS: agent in grid, picking up and dropping pkgs

For **proving** $w(G) \leq k$:

- **Necessary 1:** If a_1, \dots, a_n is optimal plan for goal G , each **prefix** a_1, \dots, a_i must be optimal plan for some t_i , $|t_i| \leq k$
- **Necessary 2:** For these t_i 's, **all** optimal plans for t_i **extend** into optimal plans for t_{i+1} .

Part II: Languages

- Language for expressing **dynamics**
- Language for expressing **general policies**
- Language for expressing **general subgoal structures**

Language for Expressing Dynamics: First-Order STRIPS

Problems specified as **instances** $P = \langle D, I \rangle$ of **general** planning domain:

- **Domain** D specified in terms of **action schemas** and **predicates**
- **Instance** is $P = \langle D, I \rangle$ where I details **objects, init, goal**

Distinction between **general** domain D and **specific** instance $P = \langle D, I \rangle$ important for **reusing** action models, and also for **learning** them:

- Learning $P_i = \langle D, I_i \rangle$ implies learning D that **generalizes** to other instances

In RL and DRL, there is no notion of **domain**: generalization to other “instances” analyzed **experimentally**; closest things are “procedurally generated instances,” and “probability distribution over tasks”

Example: 2-Gripper Problem $P = \langle D, I \rangle$ in PDDL

```
(define (domain gripper)
  (:requirements :typing)
  (:types room ball gripper)
  (:constants left right - gripper)
  (:predicates (at-robot ?r - room)(at ?b - ball ?r - room)(free ?g - gripper)
    (carry ?o - ball ?g - gripper))
  (:action move
    :parameters (?from ?to - room)
    :precondition (at-robot ?from)
    :effect (and (at-robot ?to) (not (at-robot ?from))))
  (:action pick
    :parameters (?obj - ball ?room - room ?gripper - gripper)
    :precondition (and (at ?obj ?room) (at-robot ?room) (free ?gripper))
    :effect (and (carry ?obj ?gripper) (not (at ?obj ?room)) (not (free ?gripper))))
  (:action drop
    :parameters (?obj - ball ?room - room ?gripper - gripper)
    :precondition (and (carry ?obj ?gripper) (at-robot ?room))
    :effect (and (at ?obj ?room) (free ?gripper) (not (carry ?obj ?gripper)))))

(define (problem gripper2)
  (:domain gripper)
  (:objects roomA roomB - room Ball1 Ball2 - ball)
  (:init (at-robot roomA) (free left) (free right) (at Ball1 roomA)(at Ball2 roomA))
  (:goal (and (at Ball1 roomB) (at Ball2 roomB))))
```

Preview: Learning Dynamics in Lifted STRIPS

- Planning problem $P_i = \langle D, I_i \rangle$ defines unique **state graph** $G(P_i)$
- Learning as **inverse problem**: from graphs G_1, \dots, G_k , learn D, I_i :

Given graphs G_1, \dots, G_k , find **simplest** instances $P_i = \langle D, I_i \rangle$ such that graphs G_i and $G(P_i)$ are isomorphic, $i = 1, \dots, k$.

- Problem cast and solved as combinatorial optimization task [B. and G., 2020]
- **Complexity** of D determined by $\#$ and arities of action schemas and predicates
- **Variations**: missing edges, noisy observations [Rodriguez *et al.*, 2021a]
- **Related**
 - ▷ Learning schemas from **ground traces** [Cresswell *et al.*, 2013]
 - ▷ Deep learning of action schemas from images via **autoencoders** [Asai, 2019]
 - ▷ Learning prop. action models from **options** [Konidaris *et al.*, 2018]
 - ▷ Most work on **learning action models** assumes **domain predicates** known

Second Task: General Policies

- **General policy** represents strategy for solving **multiple** domain instances **reactively**; i.e., without having to search or plan
 - ▷ E.g., policy for achieving $on(x, y)$; **any** # of blocks, **any** configuration
- What are good **languages** for expressing such policies?
- Number of works have addressed the problem [Khardon 1999; Martin and G., 2004; Fern *et al.*, 2006; Srivastava *et al.*, 2011; Hu and De Giacomo, 2011]
- **Subtlety**: set of (ground) actions change from instance to instance with objects

Learning general policies also a key goal in (Deep) RL

General Policies: A Language [B. and G., 2018]

- **General policies** are given by **rules** $C \mapsto E$ over set Φ of **features**
- **Features** f are state functions that have a well-defined value $f(s)$ on every reachable state of any instance of the domain
 - ▷ **Boolean** features p : $p(s)$ is true or false
 - ▷ **Numerical** features n : $n(s)$ is non-negative integer

Computation of feature values assumed to be “cheap”: features assumed to have **linear** number of values at most, computable in **linear** time (in $|P|$).

Example: General Policy for $clear(x)$

- **Features** $\Phi = \{H, n\}$: 'holding' and 'number of blocks above x '
- **Policy** π for class \mathcal{Q} of Block problems with goal $clear(x)$ given by two rules:

$$\{\neg H, n > 0\} \mapsto \{H, n\downarrow\} \quad ; \quad \{H, n > 0\} \mapsto \{\neg H\}$$

Meaning:

- if $\neg H$ & $n > 0$, move to successor state where H holds and n **decreases**
- if H & $n > 0$, move to successor state where $\neg H$ holds, n **doesn't change**

Language and Semantics of General Policies: Definitions

- **Policy rules** $C \mapsto E$ over set Φ of Boolean and numerical **features** p, n :
 - ▷ *Boolean conditions* in C : $p, \neg p, n = 0, n > 0$
 - ▷ *qualitative effects* in E : $p, \neg p, p?, n\downarrow, n\uparrow, n?$
- **State transition** (s, s') **satisfies** rule $C \mapsto E$ if
 - ▷ $f(s)$ makes body C true
 - ▷ change from $f(s)$ to $f(s')$ satisfies E
- A **policy** π for class \mathcal{Q} of problems P is given by policy rules $C \mapsto E$
 - ▷ *Transition* (s, s') in P compatible with π if (s, s') satisfies a policy rule
 - ▷ *Trajectory* s_0, s_1, \dots compatible if s_0 of P and transitions compatible with π
- π **solves** P if all max trajectories compatible with π reach goal of P
- π **solves** collection of problems \mathcal{Q} if it solves each $P \in \mathcal{Q}$

Example: Delivery

- Pick packages spread in $n \times m$ grid, one by one, to target location
- **Features** $\Phi = \{H, p, t, n\}$: hold, dist. to nearest pkg & target, # undelivered
- Policy π that solves class \mathcal{Q}_D : **any** # of pkgs and distribution, **any** grid size

$\{\neg H, p > 0\} \mapsto \{p\downarrow, t?\}$	go to nearest package
$\{\neg H, p = 0\} \mapsto \{H, p?\}$	pick it up
$\{H, t > 0\} \mapsto \{t\downarrow, p?\}$	go to target cell
$\{H, t = 0\} \mapsto \{\neg H, n\downarrow, p?\}$	drop package

General Policies: Three Questions

1. How to **prove** that general policy solves potentially infinite class of instances Q ?
2. How to **learn** policies (and the features involved) to solve Q ?
3. How to **learn** policies that are **guaranteed** to solve infinite Q ?

We consider idea of **learning** first and move then to 1. Not much to say about 3.

Preview: Learning General Policies

Given a known domain D , training instances P_1, \dots, P_k , over D , and a **finite pool of domain features** \mathcal{F} , each with a cost, find the cheapest policy π over \mathcal{F} such that π solves all P_i , $i = 1, \dots, k$

- Problem cast and solved as **combinatorial opt. task** [Francès *et al.*, 2021]
- Pool of **features** \mathcal{F} generated from domain predicates using **2-variable** (description) logic grammar; feature cost given by syntax tree size
- **Deep learning** approaches [Toyer *et al.*, 2018; Garg *et al.*, 2020] do not need \mathcal{F} but not 100% correct in general
- Recent DL approach also avoids \mathcal{F} and nearly 100% correct when **2-variable logic** features suffice; exploits relation between **GNNs** and 2-variable logic [Ståhlberg *et al.*, 2022a and 2022b]

Proving that a General Policy Solves Class of Instances \mathcal{Q}

How to **prove** that this policy π achieves $clear(x)$ in all Block problems?

$$\{\neg H, n > 0\} \mapsto \{H, n \downarrow\} \quad ; \quad \{H, n > 0\} \mapsto \{\neg H\}$$

- **Soundness:** policy π applies in every **non-goal** state s
 - ▷ for any such s , there is (s, s') compatible with π
- **Acyclicity:** no sequence of transitions (s_i, s_{i+1}) compatible with π **cycle**

Theorem: If π is sound and acyclic in \mathcal{Q} , and no dead-ends, π solves \mathcal{Q}

Exercise: Show that policy for $clear(x)$ is **sound** and **acyclic** in *Blocks*

Acyclicity, Termination, and QNPs

- **Termination:** criterion that ensures that policy is **acyclic** over **any** domain
- A policy π is **terminating** if for all infinite trajectories s_0, \dots, s_i, \dots compatible with π , there is a **numerical feature** n such that:
 - ▷ n is **decremented** in some recurrent transition (s, s') ; i.e., $n(s') < n(s)$
 - ▷ n is **not incremented** in any recurrent transition (s, s') ; i.e., $n(s') \not> n(s)$
- Every such trajectory deemed **impossible** or **unfair** (n can't decrement below 0), thus if π terminates, π -trajectories **terminate**
- **Termination** notion is from **QNPs**; verifiable in time $O(2^{|\Phi|})$ by SIEVE algorithm [Srivastava *et al.*, 2011], where Φ is set of features involved in the policy

More about QNPs later on . . .

Third Task: Subgoal Structure

Subgoal structure important in planning and RL (“intrinsic rewards”, hierarchies)

Sketches powerful language for expressing subgoal structure [B. and G., 2021]

- **Goal serialization** and **full policies** expressible as sketches
- **Semantics** in terms of **subgoals to be achieved**; not so with HTNs, LTL
- Sketches **split** problems into **subproblems**

If **subproblems** have a **bounded width**, problems solved in **polytime**

Example: Sketches for Delivery

- **Width=0** Sketch (full policy)

$\{\neg H, p > 0\} \mapsto \{p\downarrow, t?\}$	go to nearest package
$\{\neg H, p = 0\} \mapsto \{H, p?\}$	pick it up
$\{H, t > 0\} \mapsto \{t\downarrow, p?\}$	go to target cell
$\{H, t = 0\} \mapsto \{\neg H, n\downarrow, p?\}$	drop package

- **Width=2** Sketch:

$\{n > 0\} \mapsto \{n\downarrow\}$	deliver package
-------------------------------------	-----------------

- **Width=1** Sketch:

$\{\neg H\} \mapsto \{H\}$	go and pick package
$\{H\} \mapsto \{\neg H, n\downarrow\}$	go and deliver package

Features: Holding (H); Dist. to nearest Pkg (p), Target (t); # Undeliv Pkgs (n)

Syntax and Semantics of Sketch Rules

- **Syntax:** For Boolean and numerical **features** p and n :
 - ▷ $p, \neg p, n > 0, n = 0$ can appear in C
 - ▷ $p, \neg p, n\uparrow, n\downarrow, n?$ can appear in E
- **Semantics:** **State pair** (s, s') **satisfies** sketch rule $C \mapsto E$ if
 - ▷ $f(s)$ satisfies C
 - ▷ $(f(s), f(s'))$ satisfies E

Syntax of sketches and policies the **same**, and so with semantics, **except** that (s, s') is not a **1-step state transition** necessarily

Interpretation: When in state s , the **set of subgoal states** $G_R(s)$ to aim at is:

$$G_R(s) = \{ s' \mid (s, s') \text{ satisfies sketch rule or } s' \text{ is goal} \}$$

Sketch Width

- Sketch R **splits** problems P in \mathcal{Q} into collection of **subproblems** $P[s, G_R(s)]$:
 - ▷ **Initial state** s : reachable state s in P
 - ▷ **(Sub) goal states** $G_R(s) = \{ s' \mid (s, s') \text{ satisfies sketch rule or } s' \text{ is goal} \}$
- **Width of sketch** R over $\mathcal{Q} = \max_{s, P \in \mathcal{Q}} \text{width}(P[s, G_R(s)])$
 - ▷ *for definition in presence of **dead-ends**, see refs*

Theorem: Any P in \mathcal{Q} is **solvable** in $O(b \cdot N^{|\Phi|+2k-1})$ time by SIW_R algorithm if sketch R is **terminating** and has **width** over \mathcal{Q} bounded by k [B. and G., 2021]

▷ N : Number of atoms in problem P ; Φ : Set of features in sketch

SIW_R is like SIW but **subgoal** to achieve next given by sketch

▷ SIW is SIW_R with sketch R with **single rule**: $\{\#g > 0\} \mapsto \{\#g \downarrow\}$

Another Example: IPC Grid [Drexler et al., 2021]

This sketch is **terminating** and has **width** 1 for IPC domain Grid (pick and deliver keys spread in grid where cells can be locked and opened with other keys):

- **Sketch:**

- ▷ $r_1 : \{l > 0\} \mapsto \{l\downarrow, k?, o?, t?\}$ (if locked cells, unlock them)
- ▷ $r_2 : \{l = 0, k > 0\} \mapsto \{k\downarrow, o?, t?\}$ (else, place keys in targets)
- ▷ $r_3 : \{l > 0, \neg o\} \mapsto \{o, t?\}$ (if locked cells, pick key to open locked cell)
- ▷ $r_4 : \{l = 0, \neg t\} \mapsto \{o?, t\}$ (if all locks open and misplaced keys, pick up such key)

- **Features:**

- ▷ l is the number of unlocked grid cells
- ▷ k is the number of misplaced keys
- ▷ o is true iff robot holds key for which there is a closed lock
- ▷ t is true iff robot holds key that must be placed at some target grid cell

Preview: Learning Sketches [Drexler et al., 2022]

Given a known domain D , training instances P_1, \dots, P_n , and non-negative integer k , find simplest sketch R over a pool of features \mathcal{F} such that

- Subproblems induced by R on each P_i have all **width bounded** by k ,
- Sketch R is **terminating**

Possibly first approach for **learning subgoal structure** based on crisp **principles**

Many threads that come together:

- Planning **width**
- Language of **general policies**
- Termination notion from **QNP**s
- Semantics of **sketches**

Exercise: Test Your Knowledge! (Not trivial)

In the 1985 AIJ paper, *Macro-Operators: A Weak Method for Learning*, Rich Korf provides **macro-tables** for puzzles like Rubik Cube, 8-puzzle, and other hard puzzles that encode **policies** $\pi(s)$ for solving them from any initial state

- Can these compact policies be replaced by even more compact **sketches** of **bounded width**?
- Can these sketches be **general**? That is, applicable to Rubik cubes and n -sliding puzzles of **different sizes**?
- Can such sketches be **learned** with current method? Expressivity? Scalability? Other methods?

Background 2:

Qualitative Numerical Planning Problems (QNPs)

Language for QNPs

- Language for planning involving **propositional** and **numerical variables**
- QNPs [Srivastava *et al.* 2011] different than **numerical planning**:
 - ▷ Numerical vars in QNPs are non-negative, **real-valued**
 - ▷ **Effects** on numerical variables: just **qualitative** increments/decrements
 - ▷ **Numerical literals**: whether variable is **zero** or **positive** only
- These differences make plan-existence for QNPs **decidable**
- QNPs provide language for **general policies and sketches**:
 - ▷ QNP actions similar to policy/sketch rules but **features** replaced by **variables**
- We follow [B. and G., 2020b]

Syntax for QNPs

A **qualitative numerical problem (QNP)** is tuple $Q = \langle F, V, I, O, G \rangle$:

- F and V are sets of propositional and numerical **variables** (not features!)
- I and G denote initial and goal states
- O : actions a with precs, and prop. and numeric effects $Pre(a)$, $Eff(a)$, $N(a)$:
 - ▷ F -literals may appear in I , G , $Pre(a)$ and $Eff(a)$
 - ▷ V -literals may appear in I , G and $Pre(a)$
 - ▷ $N(a)$ can only have expressions of the form $X\uparrow$ and $X\downarrow$ for var X in V
- V -literal is either $X = 0$ or $X > 0$ for variable X in V
- **Example:** QNP $Q_{clear} = \langle \{H\}, \{n\}, I, O, G \rangle$
 - ▷ $I = \{n > 0, \neg H\}$
 - ▷ $G = \{n = 0\}$
 - ▷ $O = \{a, b\}$ where $a = \{\neg H, n > 0\} \mapsto \{H, n\downarrow\}$ and $b = \{H\} \mapsto \{\neg H\}$
- QNP actions like policy rules above but H and n not features but **variables**

Semantics and Solutions of QNPs

- Policy π for a QNP is partial map from state s into actions such that:
 - ▷ $\pi(s) = \pi(s')$ if s and s' **qualitatively similar**: same F and V true literals
- π **solves** QNP if all maximal **QNP-fair** π -trajectories reach the goal
 - ▷ **QNP fairness**: trajectory **unfair** if numerical variable **decremented** infinite number of times and **incremented** finite number of times.

Theorem [Srivastava *et al.*, 2011]: π solves QNP Q iff π is **strong cyclic solution** of the **FOND** problem $T_D(Q)$ obtained from Q that **terminates**

- $T_D(Q)$ replaces **numerical** X by **Boolean** variable “ $X > 0$ ” (“ $X = 0$ ” is negative literal)
- **Qualitative effects** $X \uparrow$ replaced by **effect** $X > 0$
- **Qualitative effects** $X \downarrow$ replaced by **non-deterministic effect** “ $X > 0 \mid X = 0$ ”
- **Strong-cyclic**: every reachable state is connected to goal state by π

Polytime reduction from QNPs to FOND, but more complex than T_D [B. and G., 2020b]

Termination, Sieve Algorithm [Srivastava et al., 2011]

Policy for QNP Q **terminates** if no infinite **QNP-fair** π -trajectories

SIEVE provides **sound** and **complete** polynomial termination test

- State s **terminates** if either
 - ▷ there is no cycle on state s , or
 - ▷ every cycle on s contains a state s' that terminates, or
 - ▷ $\pi(s)$ decrements a variable X , and every cycle on s that contains a state s' such that $\pi(s')$ increments X , contains another state s'' that terminates
- Policy π terminates iff every state reached by π terminates

Recent FOND⁺ planner handles strong FOND, strong cyclic FOND, QNPs, and hybrids by stating **fairness assumptions** explicitly [Rodriguez *et al.* 2021b]

Part III: Learning Dynamics, Policies, Sketches

- Learning **action models**:

Given graphs G_1, \dots, G_k , find **simplest** instances $P_i = \langle D, I_i \rangle$ such that graphs G_i and $G(P_i)$ are isomorphic, $i = 1, \dots, k$.

- Learning **general policies**:

Given known domain D , training instances P_1, \dots, P_k , over D , and **finite pool of domain features** \mathcal{F} , each with a cost, find the cheapest policy π over \mathcal{F} such that π solves all P_i , $i = 1, \dots, k$

- Learning **sketches**:

Given known domain D , training instances P_1, \dots, P_n , and non-negative integer k , find simplest sketch R over a pool of features \mathcal{F} such that

- ▷ Subproblems induced by R on each P_i have all **width bounded** by k ,
- ▷ Sketch R is **terminating**

Learning Action Models: Encoding [Rodriguez et al., 2021a]

- Construct **answer set program**, bounding number of objects, preds, and action/pred. arities:
 - ▷ **Given** G_1, \dots, G_k as input graphs over **black-box states**, with edge labels,
 - ▷ **Check** whether there is STRIPS model D and instances I_1, \dots, I_k such that graphs $G(P_i)$ and G_i are **isomorphic**, $i = 1, \dots, k$, where $P_i = \langle D, I_i \rangle$
 - ▷ **Optimize**: sum of action and predicate arities, . . .
- **Choice variables** in program:
 - ▷ Lifted preds/effects for each action schema (schemas determined by labels in input graphs)
 - ▷ Values of ground atoms at each state
 - ▷ Assignment of applicable grounded actions to edges in input graphs
- **Constraints** in program:
 - ▷ Different nodes in each input graph maps to different valuations of grounded atoms
 - ▷ Every edge in input graph “receives” a grounded action (establishing isomorphism)
 - ▷ Compliance of preds/effects of assigned grounded actions to edges
- CLINGO program \sim 400 lines [Rodriguez et al. 2021a]; more complex in SAT [B. and G., 2020a]

Learning General Policies: Encoding [Francès et al., 2021]

- **Input** is set of transitions \mathcal{S} from small instances, pool of features \mathcal{F} , integer δ
- **Output** is policy: rules obtained from **selected features** and (“good”) **transitions**
- **Combinatorial opt. task** $T(\mathcal{S}, \mathcal{F}, \delta)$: Solve constraints minimizing **feature complexity**
 - ▷ **Choice variables:** $select(f)$, $good(s, s')$ and $V(s, d)$
 - ▷ **Constraints:**
 1. $\bigvee \{ good(s, s') : (s, s') \in \mathcal{S} \}$ (Good transition at each non-terminal s)
 2. $\neg good(s, s')$ (No good reach dead-end s')
 3. $\bigvee \{ select(f) : f \text{ such } f[s] \neq f[s'] \}$ (Distinguish $\{s, s'\}$ when exactly-1 is goal)
 4. Exactly-1 $\{ V(s, d) : V^*(s) \leq d \leq \delta V^*(s) \}$ (Set distances)
 5. $good(s, s') \wedge V(s, d) \wedge V(s', d') \rightarrow d < d'$ (Distances avoid cycles)
 6. $good(s, s') \wedge \neg good(t, t') \rightarrow D2(s, s'; t, t')$ (Distinguish good/bad transitions)

where $D2(s, s'; t, t') = \bigvee_{f: \Delta_f(s, s') \neq \Delta_f(t, t')} select(f)$ and $\Delta_f(s, s') \in \{\uparrow, \downarrow, =\}$

Learning General Sketches: Encoding [Drexler et al., 2022]

- **Input:** transitions \mathcal{S} in small instances, pool \mathcal{F} , width bound k , max # sketch rules m
- **Output:** sketch of **width** $\leq k$, **acyclic** in given instances, with up to m rules
- **Combinatorial opt. task** $T(\mathcal{S}, \mathcal{F}, k, m)$: solve constraints min **complexity** of selected features
- **Variables:**
 1. $select(f)$ (features f in pool \mathcal{F})
 2. $cond(i, f, v)$ (rule i , feature f , condition v ; e.g. \top , “=0”, etc)
 3. $eff(i, f, e)$ (rule i , feature f , effect e)
 4. $satisfies(s, s', i)$ (for s' reachable from s , rule i)
 5. $subgoal(s, t)$ (for tuple t of width $\leq k$ reachable from s)
 6. $subgoals(s, t, s')$ (some optimal plan from s to t ends in s')
- **Constraints:**
 1. $cond(i, f, v)$ and $eff(i, f, e)$ use unique v , imply $select(f)$ (well formed sketch)
 2. $\bigvee_t subgoal(s, t)$ (width k subgoal)
 3. $subgoal(s, t) \leftrightarrow subgoals(s, t, s')$ (subgoal t of s may lead to s')
 4. $subgoals(s, t, s') \rightarrow \bigvee_i satisfies(s, s', i)$ ((s, s') must satisfy some rule)
 5. $satisfies(s, s'', i) \rightarrow \bigvee \{subgoal(s, t) : t \text{ such } d(s, t) < d(s, s'')\}$ (dead-end s'' is farther)
 6. $satisfies(s, s', i) \rightarrow \bigvee \{subgoal(s, t) : t \text{ such } d(s, t) \leq d(s, s')\}$ (subgoals optimal)
 7. $satisfies(s, s', i) \leftrightarrow “(s, s') \text{ satisfies rule } i”$ (atom true if condition holds)
 8. Collection of rules is **terminating** (approx'ed by testing acyclicity)
- Paper to be presented at the conference (ICAPS 2022)

About the Pool of Features \mathcal{F} [B. et al., 2019]

- **Description logic grammar** allows generation of **concepts** and **roles** from **domain predicates**
- Complexity of concept/role given by **size of its syntax tree**
- Pool \mathcal{F} obtained from concepts of complexity bounded by parameter
- Denotation of concept C in state s is **subset of objects**
- Each concept C defines num and Bool features $n_C(s) = |C(s)|$; $p_C(s) = \top$ iff $|C(s)| > 0$
- Grammar:
 - ▷ Primitive: C_p given by unary predicates p and unary “goal predicates” p_G
 - ▷ Universal: C_u contains all objects
 - ▷ Nominals: $C_a = \{a\}$ for constants/parameter a
 - ▷ Negation: $\neg C$ contains $C_u \setminus C$
 - ▷ Intersection: $C \sqcap C'$
 - ▷ Quantified: $\exists R.C = \{x : \exists y[R(x, y) \wedge C(y)]\}$ and $\forall R.C = \{x : \forall y[R(x, y) \wedge C(y)]\}$
 - ▷ Roles (for binary predicate p): R_p , R_p^{-1} , R_p^+ , and $[R_p^{-1}]^+$
- Additional **distance features** $dist(C_1, R, C_2)$ for concepts C_1 and C_2 and role R that evaluates to d in state s iff minimum R -distance between object in C_1 to object in C_2 is d

General Policies By Deep Learning [Ståhlberg et al., 2022a,b]

- Exploits correspondence between **graph neural networks (GNNs)** and **two-variable logic \mathcal{C}_2** to learn policy **without requiring** pool of \mathcal{C}_2 features \mathcal{F}
- **Value function V** learned that yields general policy π_V **greedy** in V
- For **generalization**, based on GNN arch. for MaxCSP(Γ) [Toenshoff et al., 2021]
 - ▷ **Input** given by the states s extended with “goal predicates” p_G
 - ▷ **Output** $V(s)$ is non-linear aggregation of object embeddings
 - ▷ **Min Loss:** $|V^*(s) - V(s)|$ for supervised learning of optimal policies
 - ▷ **Min Loss:** $\max\{0, [1 + \min_{s' \in N(s)} V(s')] - V(s)\}$ unsupervised/non-optimal
- Nearly **as good as** policies based on **explicit pool \mathcal{F} of \mathcal{C}_2 features**
- Complexity of “latent features” not explicitly bounded
- Paper to be presented at the conference (ICAPS 2022)

GNN Architecture [Ståhlberg et al., 2022a,b]

Algorithm 1: GNN maps state s into scalar $V(s)$

Input: State s : set of atoms true in s , set of objects

Output: $V(s)$

```
1  $f_0(o) \sim \mathbf{0}^{k/2} \mathcal{N}(0, 1)^{k/2}$  for each object  $o \in s$ ;  
2 for  $i \in \{0, \dots, L - 1\}$  do  
3   for each atom  $q := p(o_1, \dots, o_m)$  true in  $s$  do  
4      $m_{q,o} := [\mathbf{MLP}_p(f_i(o_1), \dots, f_i(o_m))]_j$ ;  
5   for each  $o$  in  $s$  do  
6      $f_{i+1}(o) := \mathbf{MLP}_U(f_i(o), \text{agg}(\{m_{q,o} | o \in q\}))$ ;  
    $//$  Aggregate, update embeddings  
7  $V := \mathbf{MLP}_2(\sum_{o \in s} \mathbf{MLP}_1(f_L(o)))$ 
```

Wrap Up: Representation Learning for Acting and Planning

- **Background 1:** Classical planning, planning **width**
- **Languages** for
 - ▷ representing general dynamics
 - ▷ representing general policies
 - ▷ representing general subgoal structures (sketches; ‘intrinsic rewards’)
- **Background 2:** Qualitative numerical planning problems (**QNP**s)
- **Learning** representations over these languages:
 - ▷ learning general dynamics
 - ▷ learning general policies
 - ▷ learning general subgoal structures
- **Wrap up; Challenges**

Wrap Up

- To learn representations that generalize due to structure, don't play with low-level neural architecture; choose suitable (domain-independent) **target language** and learn representations over it:
 - ▷ generalization
 - ▷ transparency
 - ▷ powerful, meaningful bias
 - ▷ distinction between **what** and **how**
- Examples of learning language-based representations to **act** and **plan**:
 - ▷ general action **dynamics**
 - ▷ general **policies**
 - ▷ general **subgoal structures** (sketches)

Challenges: Language-based Representation Learning

- Scalability of combinatorial optimization approaches
- Use of deep learning (learning lifted dynamics, policies, sketches).
- Alternative target languages for learning (e.g., vs lifted STRIPS)
- Continuous domains, space, time
- Stochastic and non-deterministic domains
- States in the input: black-box, parsed images, images, videos
- Grounded vs. ungrounded representations
- Learning and reusing “skills”, hierarchies
- . . .

Plenty to do; if seriously interested, reach us

References

- [Asai, 2019] Asai, M. (2019). Unsupervised grounding of plannable first-order logic representation from images. In *Proc. ICAPS*.
- [Bonet et al., 2019] Bonet, B., Francès, G., and Geffner, H. (2019). Learning features and abstract actions for computing generalized plans. In *Proc. AAAI*, pages 2703–2710.
- [Bonet and Geffner, 2018] Bonet, B. and Geffner, H. (2018). Features, projections, and representation change for generalized planning. In *Proc. IJCAI*, pages 4667–4673.
- [Bonet and Geffner, 2020a] Bonet, B. and Geffner, H. (2020a). Learning first-order symbolic representations for planning from the structure of the state space. In *Proc. ECAI*.
- [Bonet and Geffner, 2020b] Bonet, B. and Geffner, H. (2020b). Qualitative numeric planning: Reductions and complexity. *Journal of AI Research*, 69:923–961.
- [Bonet and Geffner, 2021] Bonet, B. and Geffner, H. (2021). General policies, representations, and planning width. In *Proc. AAAI*, pages 11764–11773.
- [Chevalier-Boisvert et al., 2019] Chevalier-Boisvert, M., Bahdanau, D., Lahlou, S., Willems, L., Saharia, C., Nguyen, T. H., and Bengio, Y. (2019). Babyai: A platform to study the sample efficiency of grounded language learning. In *ICLR*.
- [Cresswell et al., 2013] Cresswell, S. N., McCluskey, T. L., and West, M. M. (2013). Acquiring planning domain models using LOCM. *The Knowledge Engineering Review*, 28(2):195–213.
- [Drexler et al., 2021] Drexler, D., Seipp, J., and Geffner, H. (2021). Expressing and exploiting the common subgoal structure of classical planning domains using sketches. In *Proc. KR*, pages 258–268.
- [Drexler et al., 2022] Drexler, D., Seipp, J., and Geffner, H. (2022). Learning sketches for decomposing planning problems into subproblems of bounded width. In *Proc. ICAPS*.
- [Fern et al., 2006] Fern, A., Yoon, S., and Givan, R. (2006). Approximate policy iteration with a policy language bias: Solving relational markov decision processes. *Journal of Artificial Intelligence Research*, 25:75–118.

- [Francès et al., 2021] Francès, G., Bonet, B., and Geffner, H. (2021). Learning general planning policies from small examples without supervision. In *Proc. AAAI*, pages 11801–11808.
- [Garg et al., 2020] Garg, S., Bajpai, A., and Mausam (2020). Symbolic network: generalized neural policies for relational mdps. In *International Conference on Machine Learning*, pages 3397–3407.
- [Geffner and Bonet, 2013] Geffner, H. and Bonet, B. (2013). *A Concise Introduction to Models and Methods for Automated Planning*. Morgan & Claypool Publishers.
- [Hu and De Giacomo, 2011] Hu, Y. and De Giacomo, G. (2011). Generalized planning: Synthesizing plans that work for multiple environments. In *Proc. IJCAI*, pages 918–923.
- [Khardon, 1999] Khardon, R. (1999). Learning action strategies for planning domains. *Artificial Intelligence*, 113:125–148.
- [Konidaris et al., 2018] Konidaris, G., Kaelbling, L. P., and Lozano-Perez, T. (2018). From skills to symbols: Learning symbolic representations for abstract high-level planning. *Journal of Artificial Intelligence Research*, 61:215–289.
- [Lipovetzky and Geffner, 2012] Lipovetzky, N. and Geffner, H. (2012). Width and serialization of classical planning problems. In *Proc. ECAI*, pages 540–545.
- [Lipovetzky and Geffner, 2017a] Lipovetzky, N. and Geffner, H. (2017a). Best-first width search: Exploration and exploitation in classical planning. In *Proc. AAAI*.
- [Lipovetzky and Geffner, 2017b] Lipovetzky, N. and Geffner, H. (2017b). A polynomial planning algorithm that beats lama and ff. *Proc. ICAPS*.
- [Martín and Geffner, 2004] Martín, M. and Geffner, H. (2004). Learning generalized policies from planning examples using concept languages. *Applied Intelligence*, 20(1):9–19.
- [Rodriguez et al., 2021a] Rodriguez, I. D., Bonet, B., Romero, J., and Geffner, H. (2021a). Learning first-order representations for planning from black-box states: New results. In *KR*. arXiv preprint arXiv:2105.10830.
- [Rodriguez et al., 2021b] Rodriguez, I. D., Bonet, B., Sardina, S., , and Geffner, H. (2021b). Flexible fond planning with explicit fairness assumptions. In *Proc. ICAPS*, pages 290–298.
- [Srivastava et al., 2011] Srivastava, S., Zilberstein, S., Immerman, N., and Geffner, H. (2011). Qualitative numeric planning. In *AAAI*.

- [Ståhlberg et al., 2022a] Ståhlberg, S., Bonet, B., and Geffner, H. (2022a). Learning general optimal policies with graph neural networks: Expressive power, transparency, and limits. In *Proc. ICAPS*.
- [Ståhlberg et al., 2022b] Ståhlberg, S., Bonet, B., and Geffner, H. (2022b). Learning generalized policies without supervision using gnns. In *Proc. KR*.
- [Toenshoff et al., 2021] Toenshoff, J., Ritzert, M., Wolf, H., and Grohe, M. (2021). Graph neural networks for maximum constraint satisfaction. *Frontiers in artificial intelligence*, 3:98.
- [Toyer et al., 2018] Toyer, S., Trevizan, F., Thiébaux, S., and Xie, L. (2018). Action schema networks: Generalised policies with deep learning. In *AAAI*.